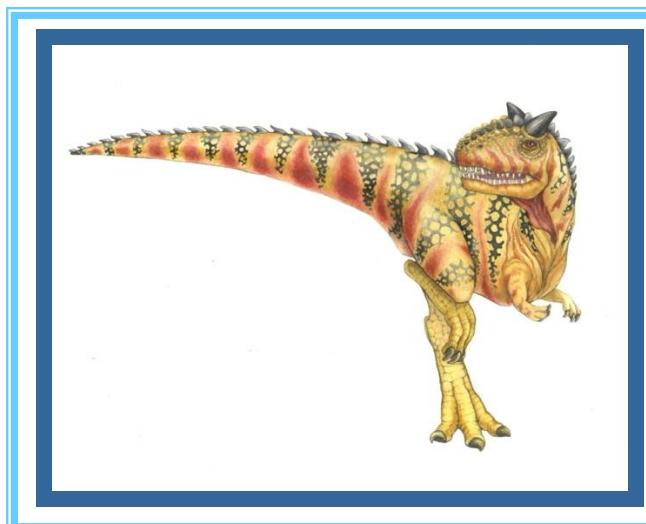
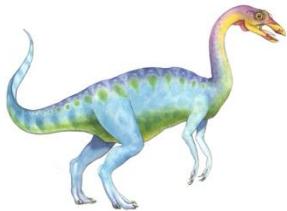


# Chapter 3: Process Concept

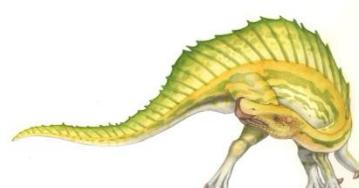




# Chapter 3: Process Concept

---

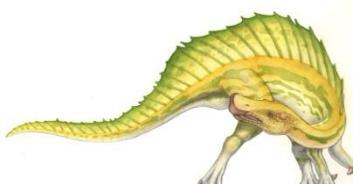
- Process Concept
- Process Scheduling
- Operations on Processes
- Interprocess Communication
- Examples of IPC Systems
- Communication in Client-Server Systems





# Objectives

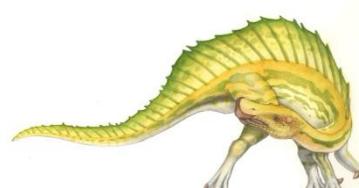
- To introduce the notion of a process -- a program in execution, which forms the basis of all computation
- To describe the various features of processes, including scheduling, creation and termination, and communication
- To explore interprocess communication using shared memory and message passing
- To describe communication in client-server systems





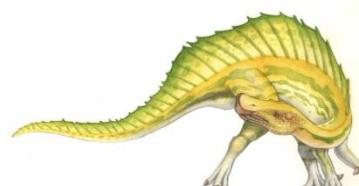
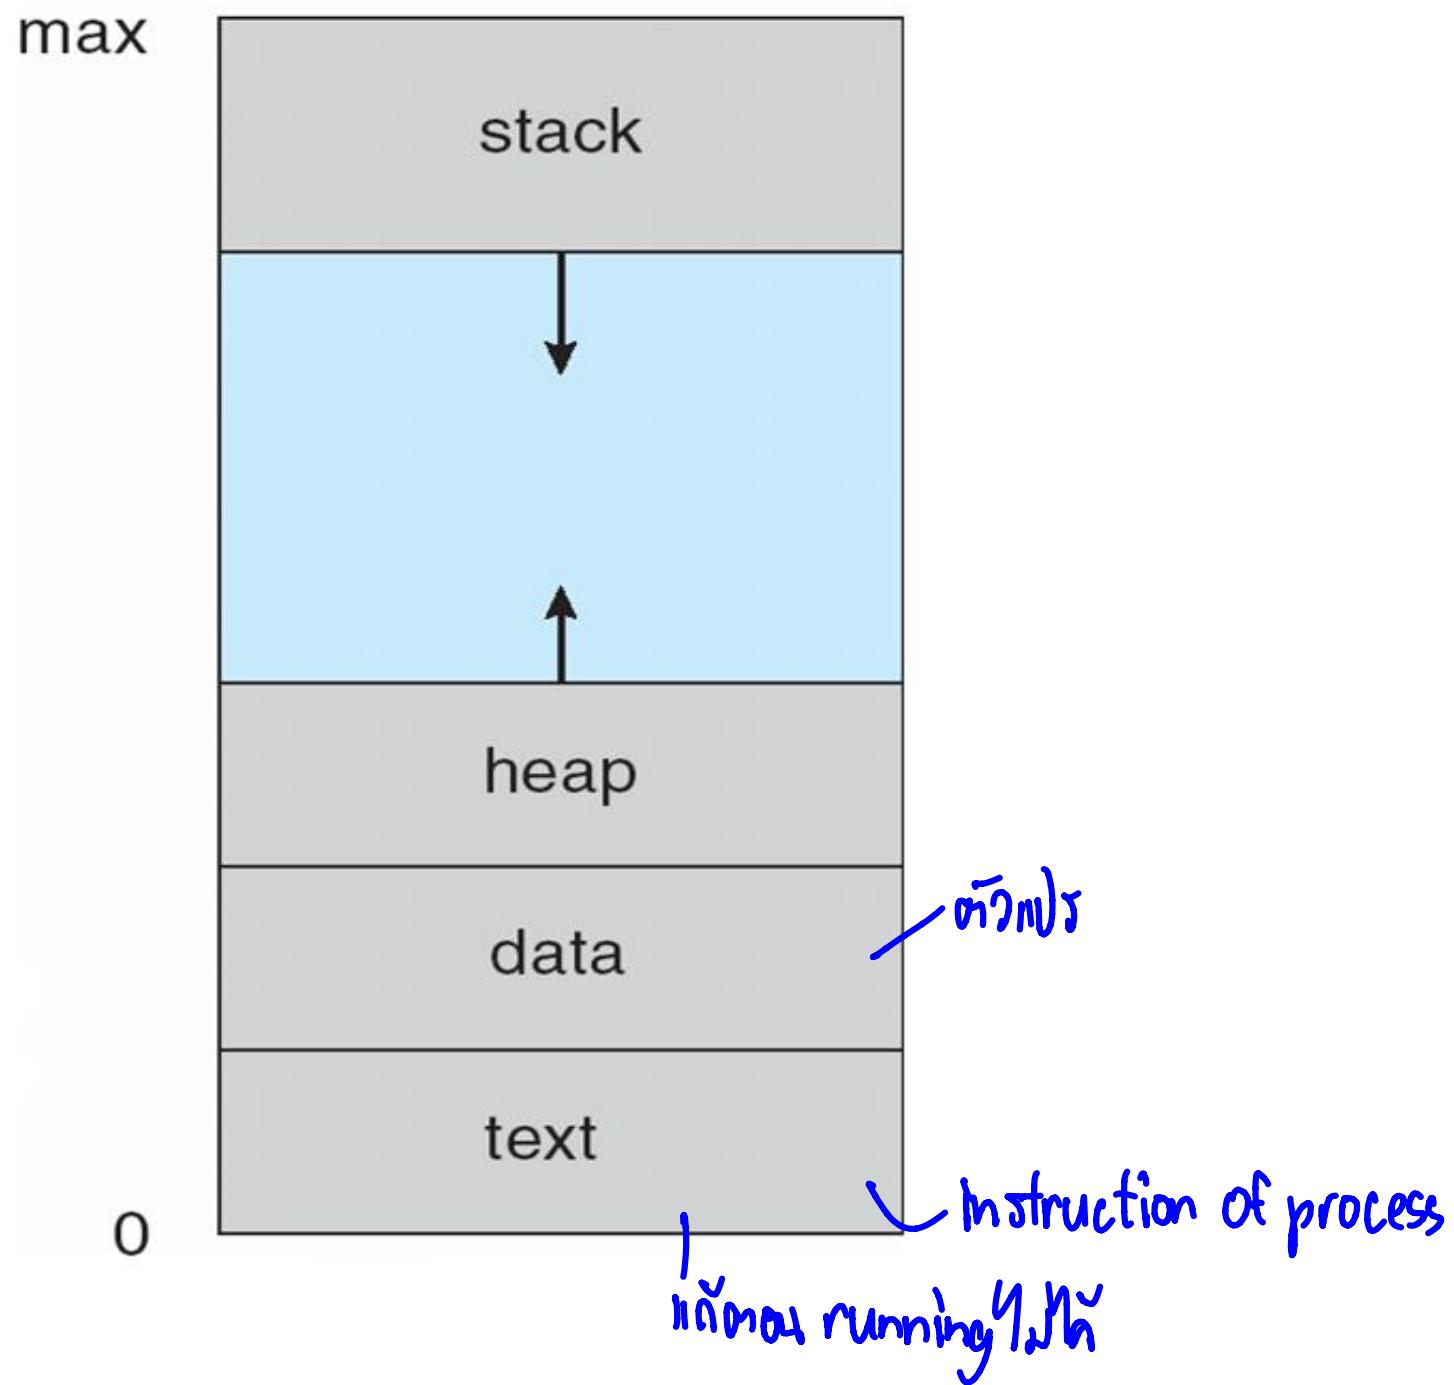
# Process Concept

- An operating system executes a variety of programs:
  - Batch system – **jobs**
  - Time-shared systems – **user programs** or **tasks**
- Textbook uses the terms **job** and **process** almost interchangeably
- **Process** – a program in execution; process execution must progress in sequential fashion
- Multiple parts
  - The **program code**, also called **text section**
  - Current activity including **program counter**, processor registers
  - **Stack** containing temporary data
    - ▶ Function parameters, return addresses, local variables
  - **Data section** containing global variables
  - **Heap** containing memory dynamically allocated during run time
- Program is **passive** entity stored on disk (**executable file**), process is **active**
  - Program becomes process when executable file loaded into memory
- Execution of program started via GUI mouse clicks, command line entry of its name, etc
- One program can be several processes
  - Consider multiple users executing the same program





# Process in Memory





# Process State

/ process life cycle

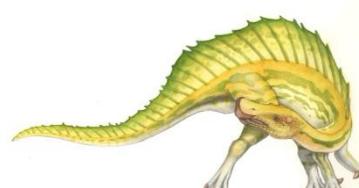
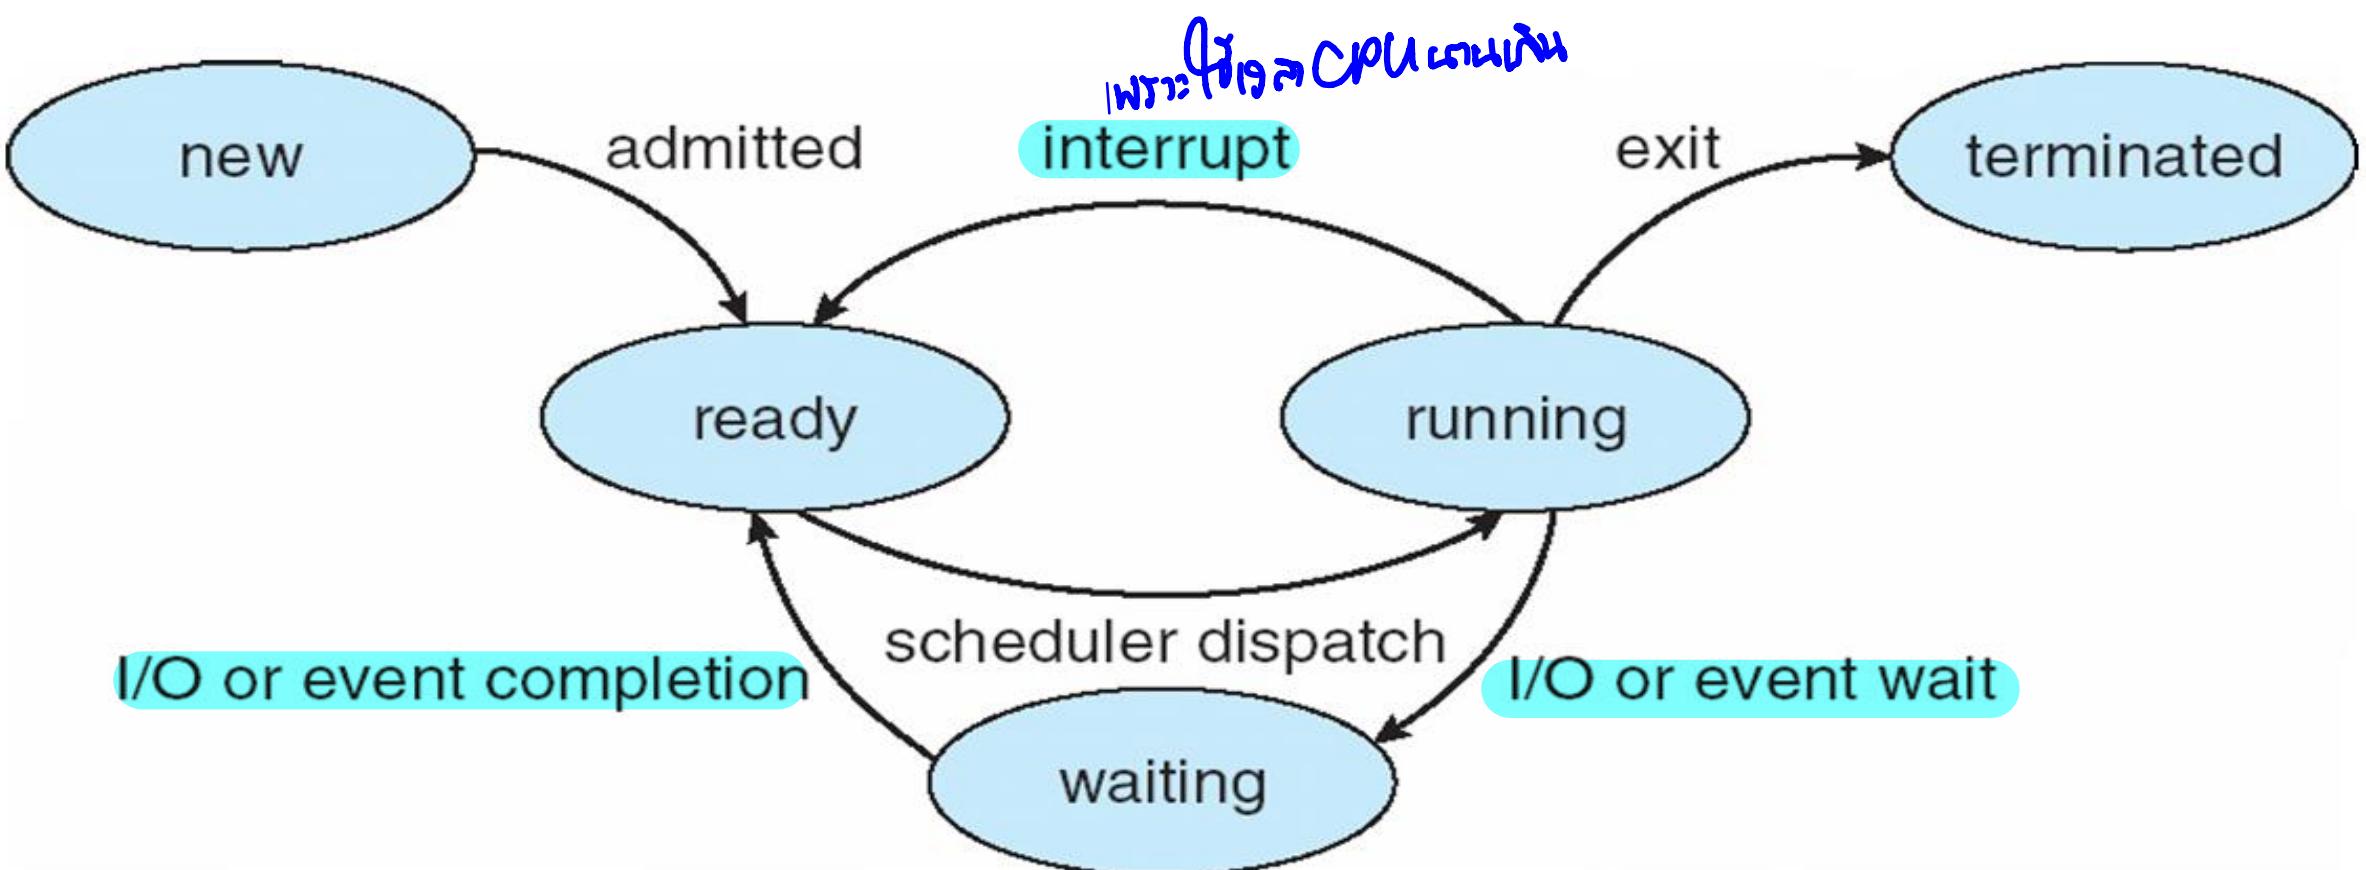
## ສຳເນົາ

- As a process executes, it changes state
- ① □ **new**: The process is being created  
→ run ຖືກສໍຕາເປັນການ #10 processor
- ③ □ **running**: Instructions are being executed ສຳເນົາ
- **waiting**<sup>ໃຫຍ່ I/O</sup>: The process is waiting for some event to occur
- ② □ **ready**: The process is waiting to be assigned to a processor <sup>ອີກ້ມາ queue</sup>
- **terminated**: The process has finished execution





# Diagram of Process State





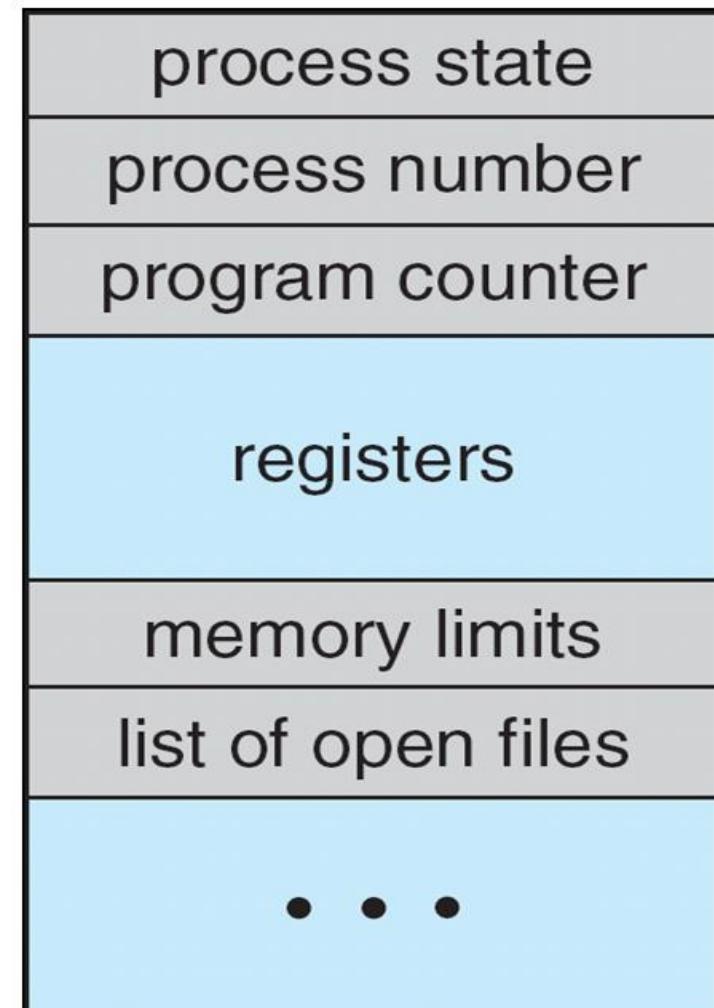
spin process

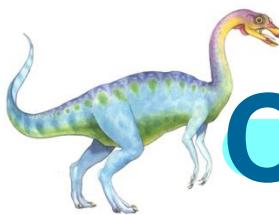
# Process Control Block (PCB)

## Information associated with each process

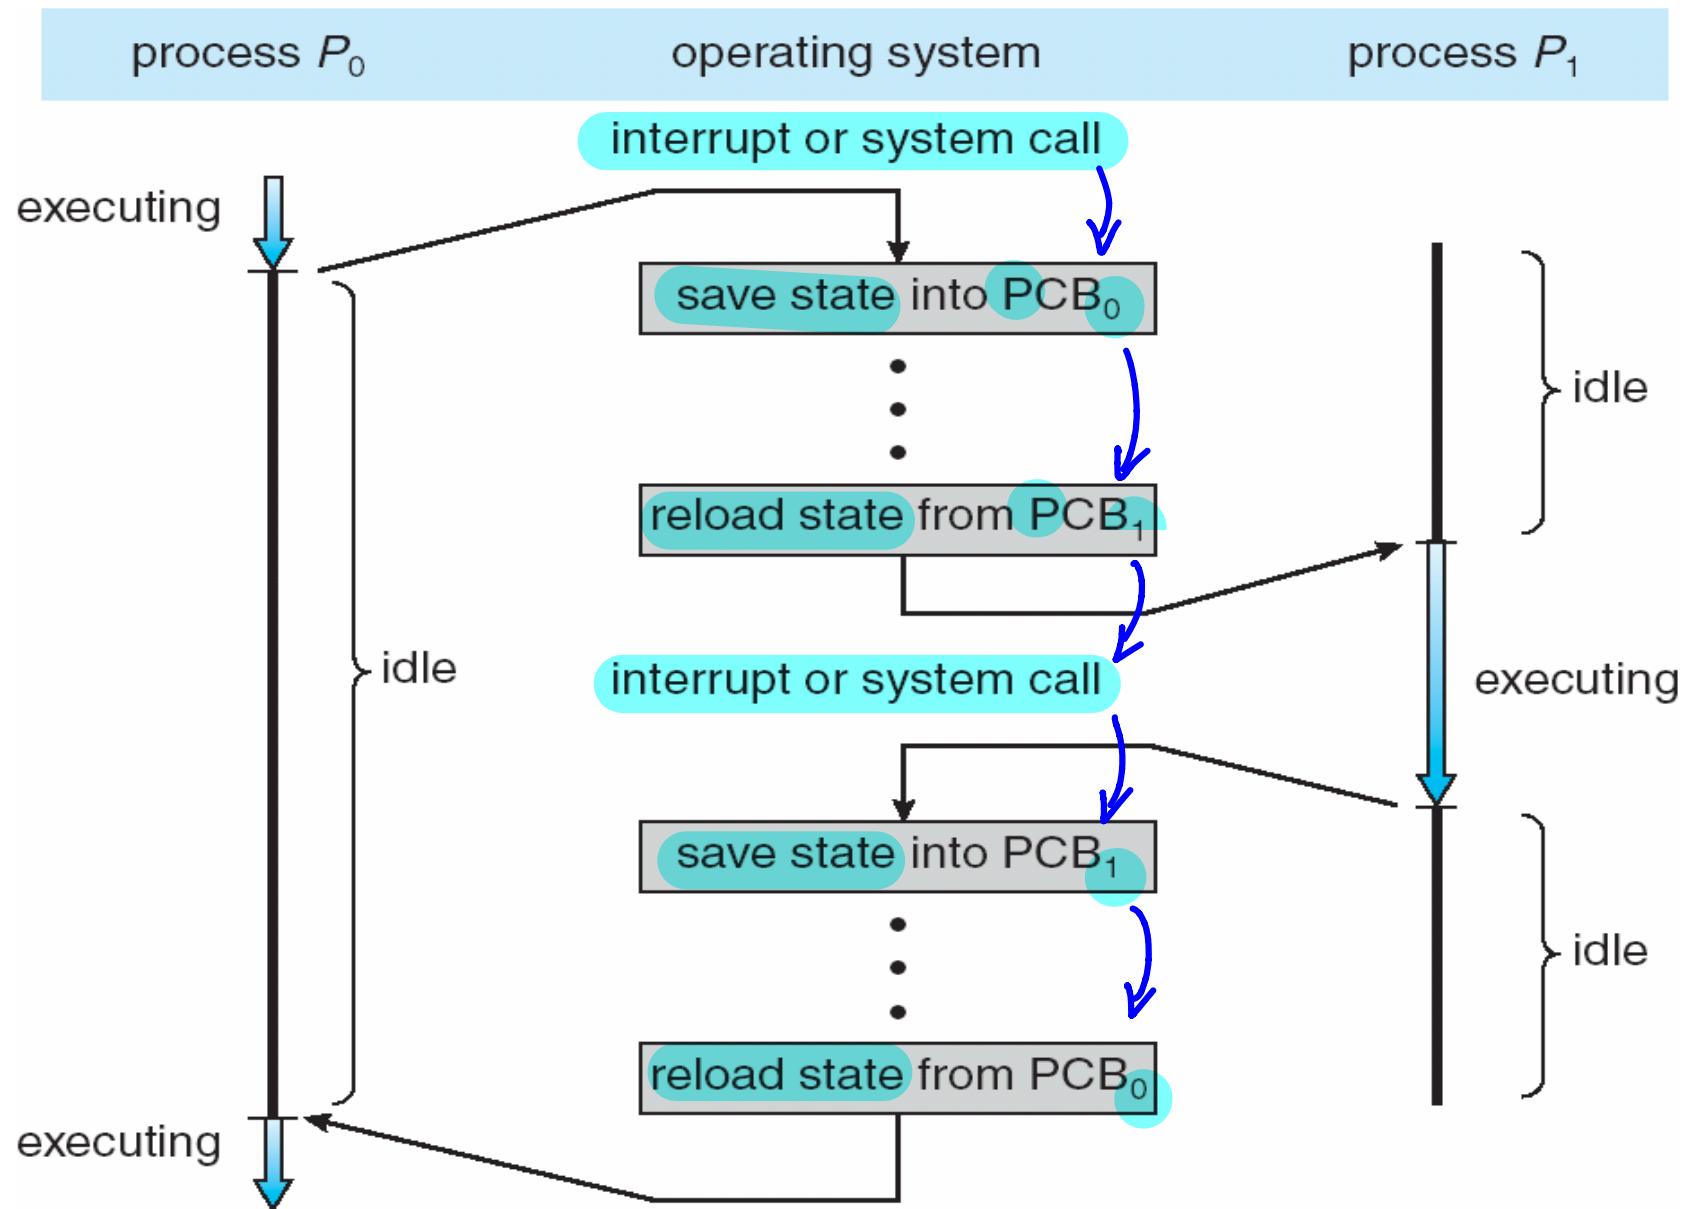
(also called **task control block**)

- ❑ Process state – running, waiting, etc
  - ❑ Program counter – location of instruction to next execute
  - ❑ CPU registers – contents of all process-centric registers
  - ❑ CPU scheduling information- priorities, scheduling queue pointers
  - ❑ Memory-management information – memory allocated to the process
  - ❑ Accounting information – CPU used, clock time elapsed since start, time limits
  - ❑ I/O status information – I/O devices allocated to process, list of open files





# CPU Switch From Process to Process



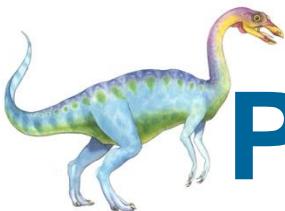


# Threads

---

- So far, process has a single thread of execution  
*downward*
- Consider having multiple program counters per process
  - Multiple locations can execute at once
    - ▶ Multiple threads of control -> **threads**
- Must then have storage for thread details, multiple program counters in PCB
- See next chapter



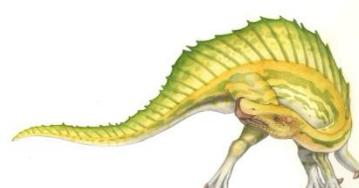
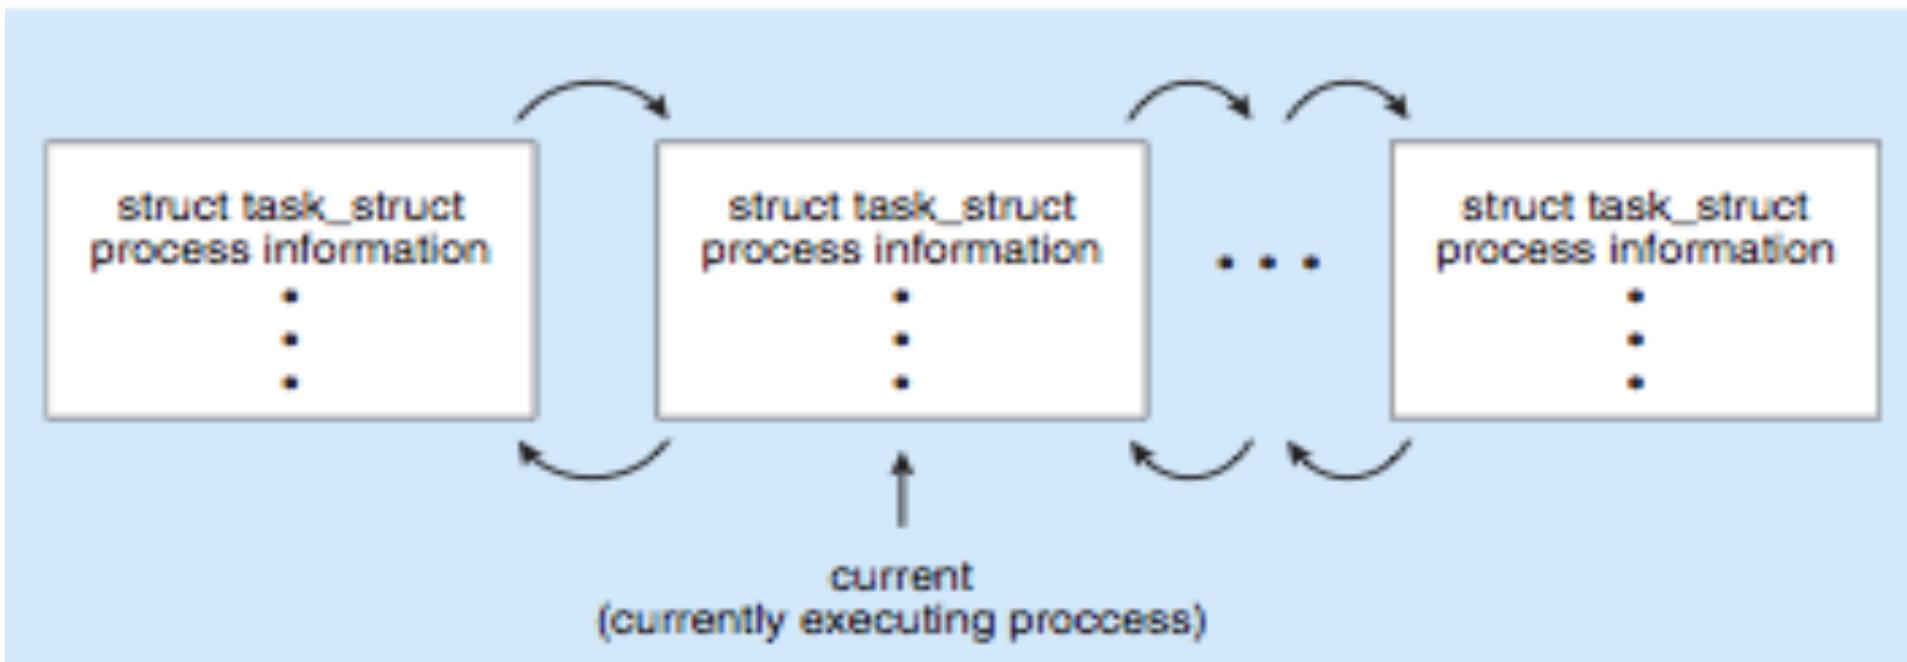


# Process Representation in Linux

## link list

- Represented by the C structure `task_struct`

```
pid_t pid; /* process identifier */  
long state; /* state of the process */  
unsigned int time_slice /* scheduling information */  
struct task_struct *parent; /* this process's parent */  
struct list_head children; /* this process's children */  
struct files_struct *files; /* list of open files */  
struct mm_struct *mm; /* address space of this process */
```

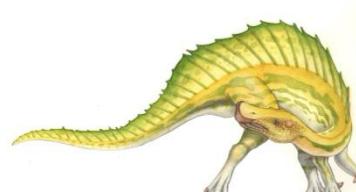




# Process Scheduling



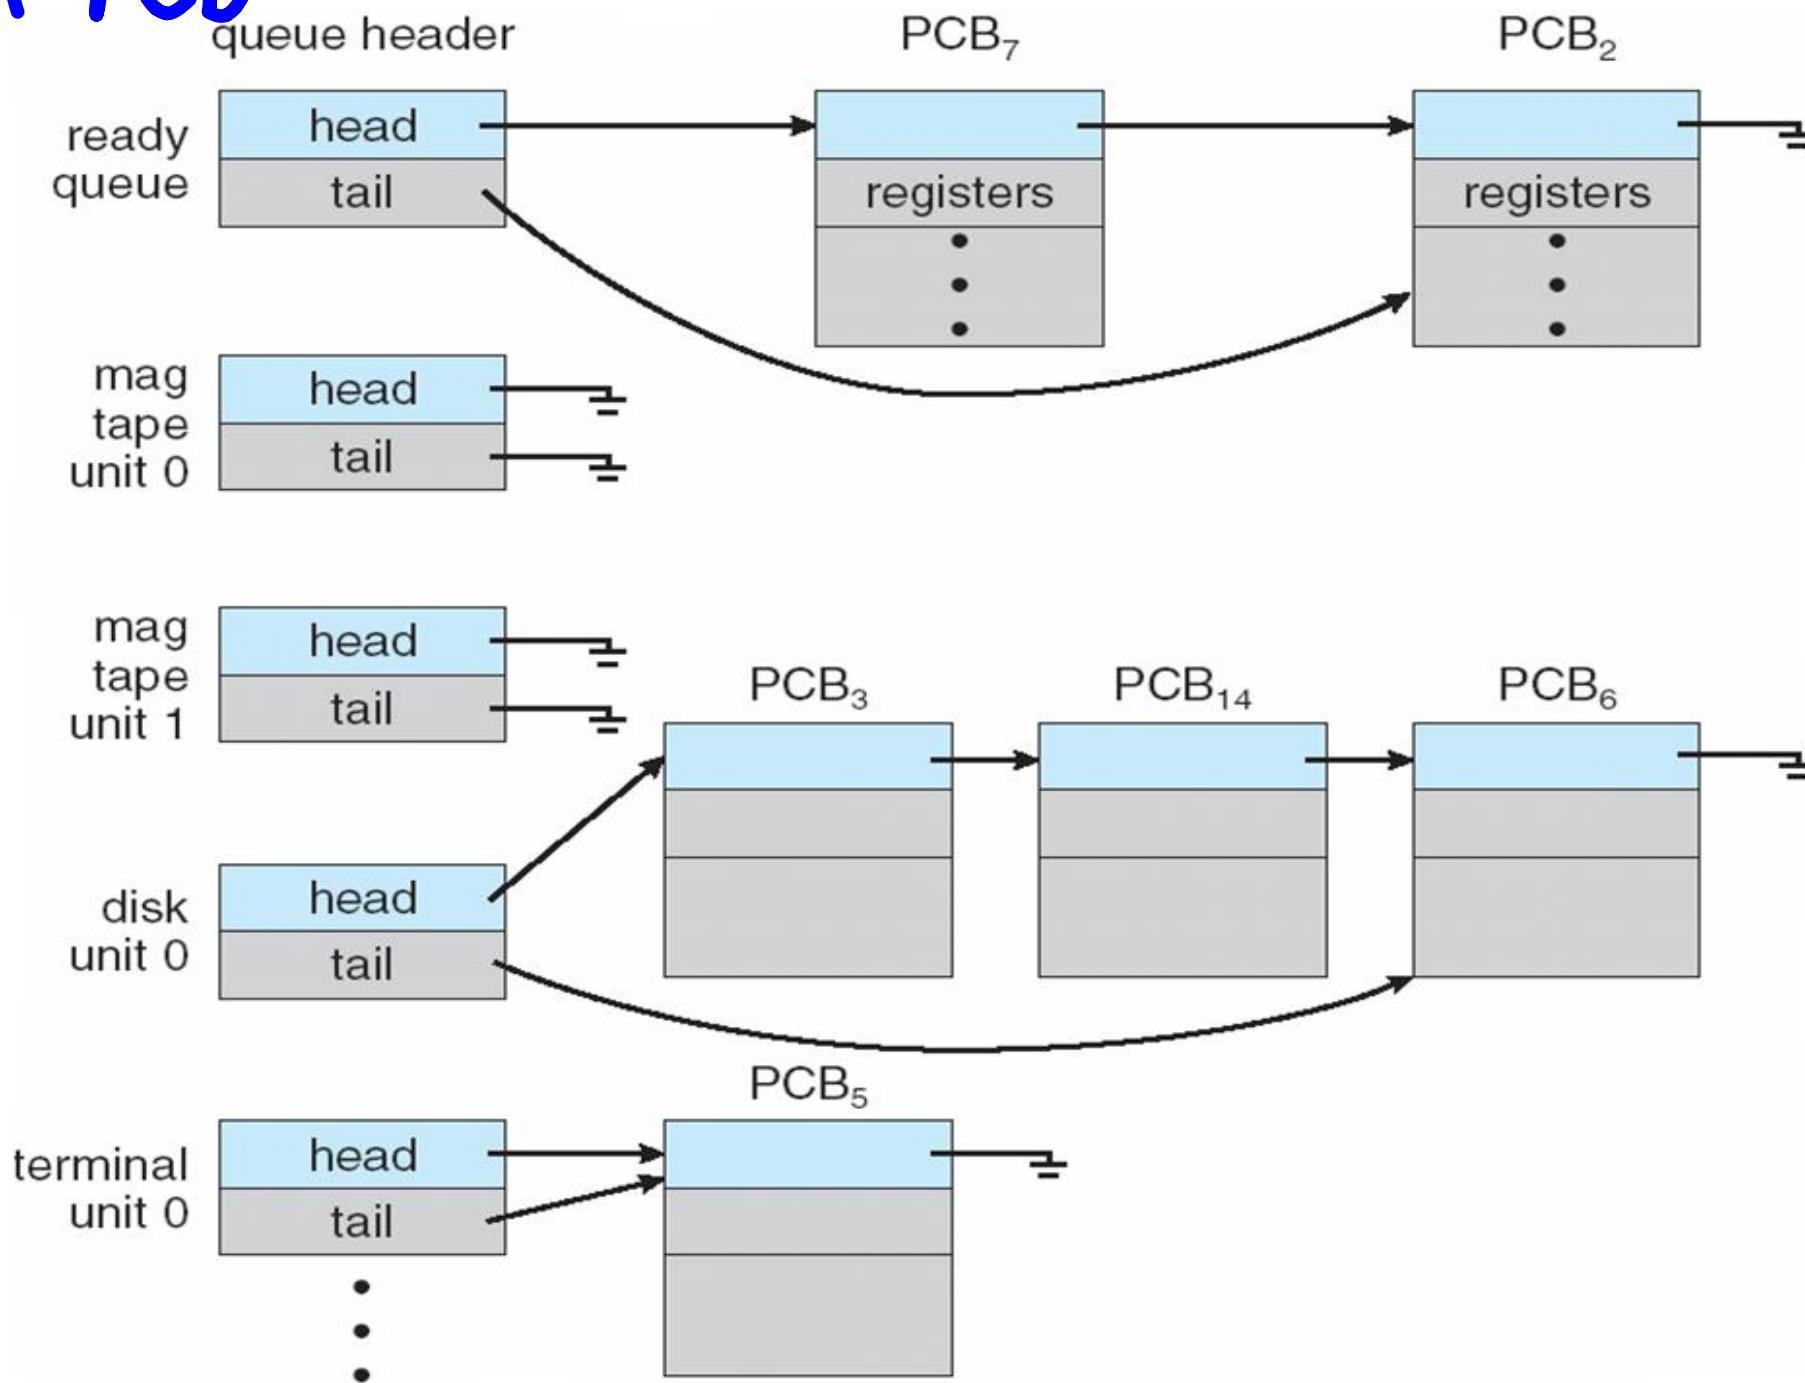
- Maximize CPU use, quickly switch processes onto CPU for time sharing
- Process scheduler selects among available processes for next execution on CPU
- Maintains scheduling queues of processes
  - Job queue – set of all processes in the system
  - Ready queue – set of all processes residing in main memory, ready and waiting to execute
  - Device queues – set of processes waiting for an I/O device
  - Processes migrate among the various queues





# Ready Queue And Various I/O Device Queues

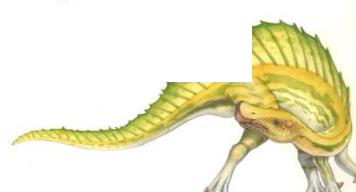
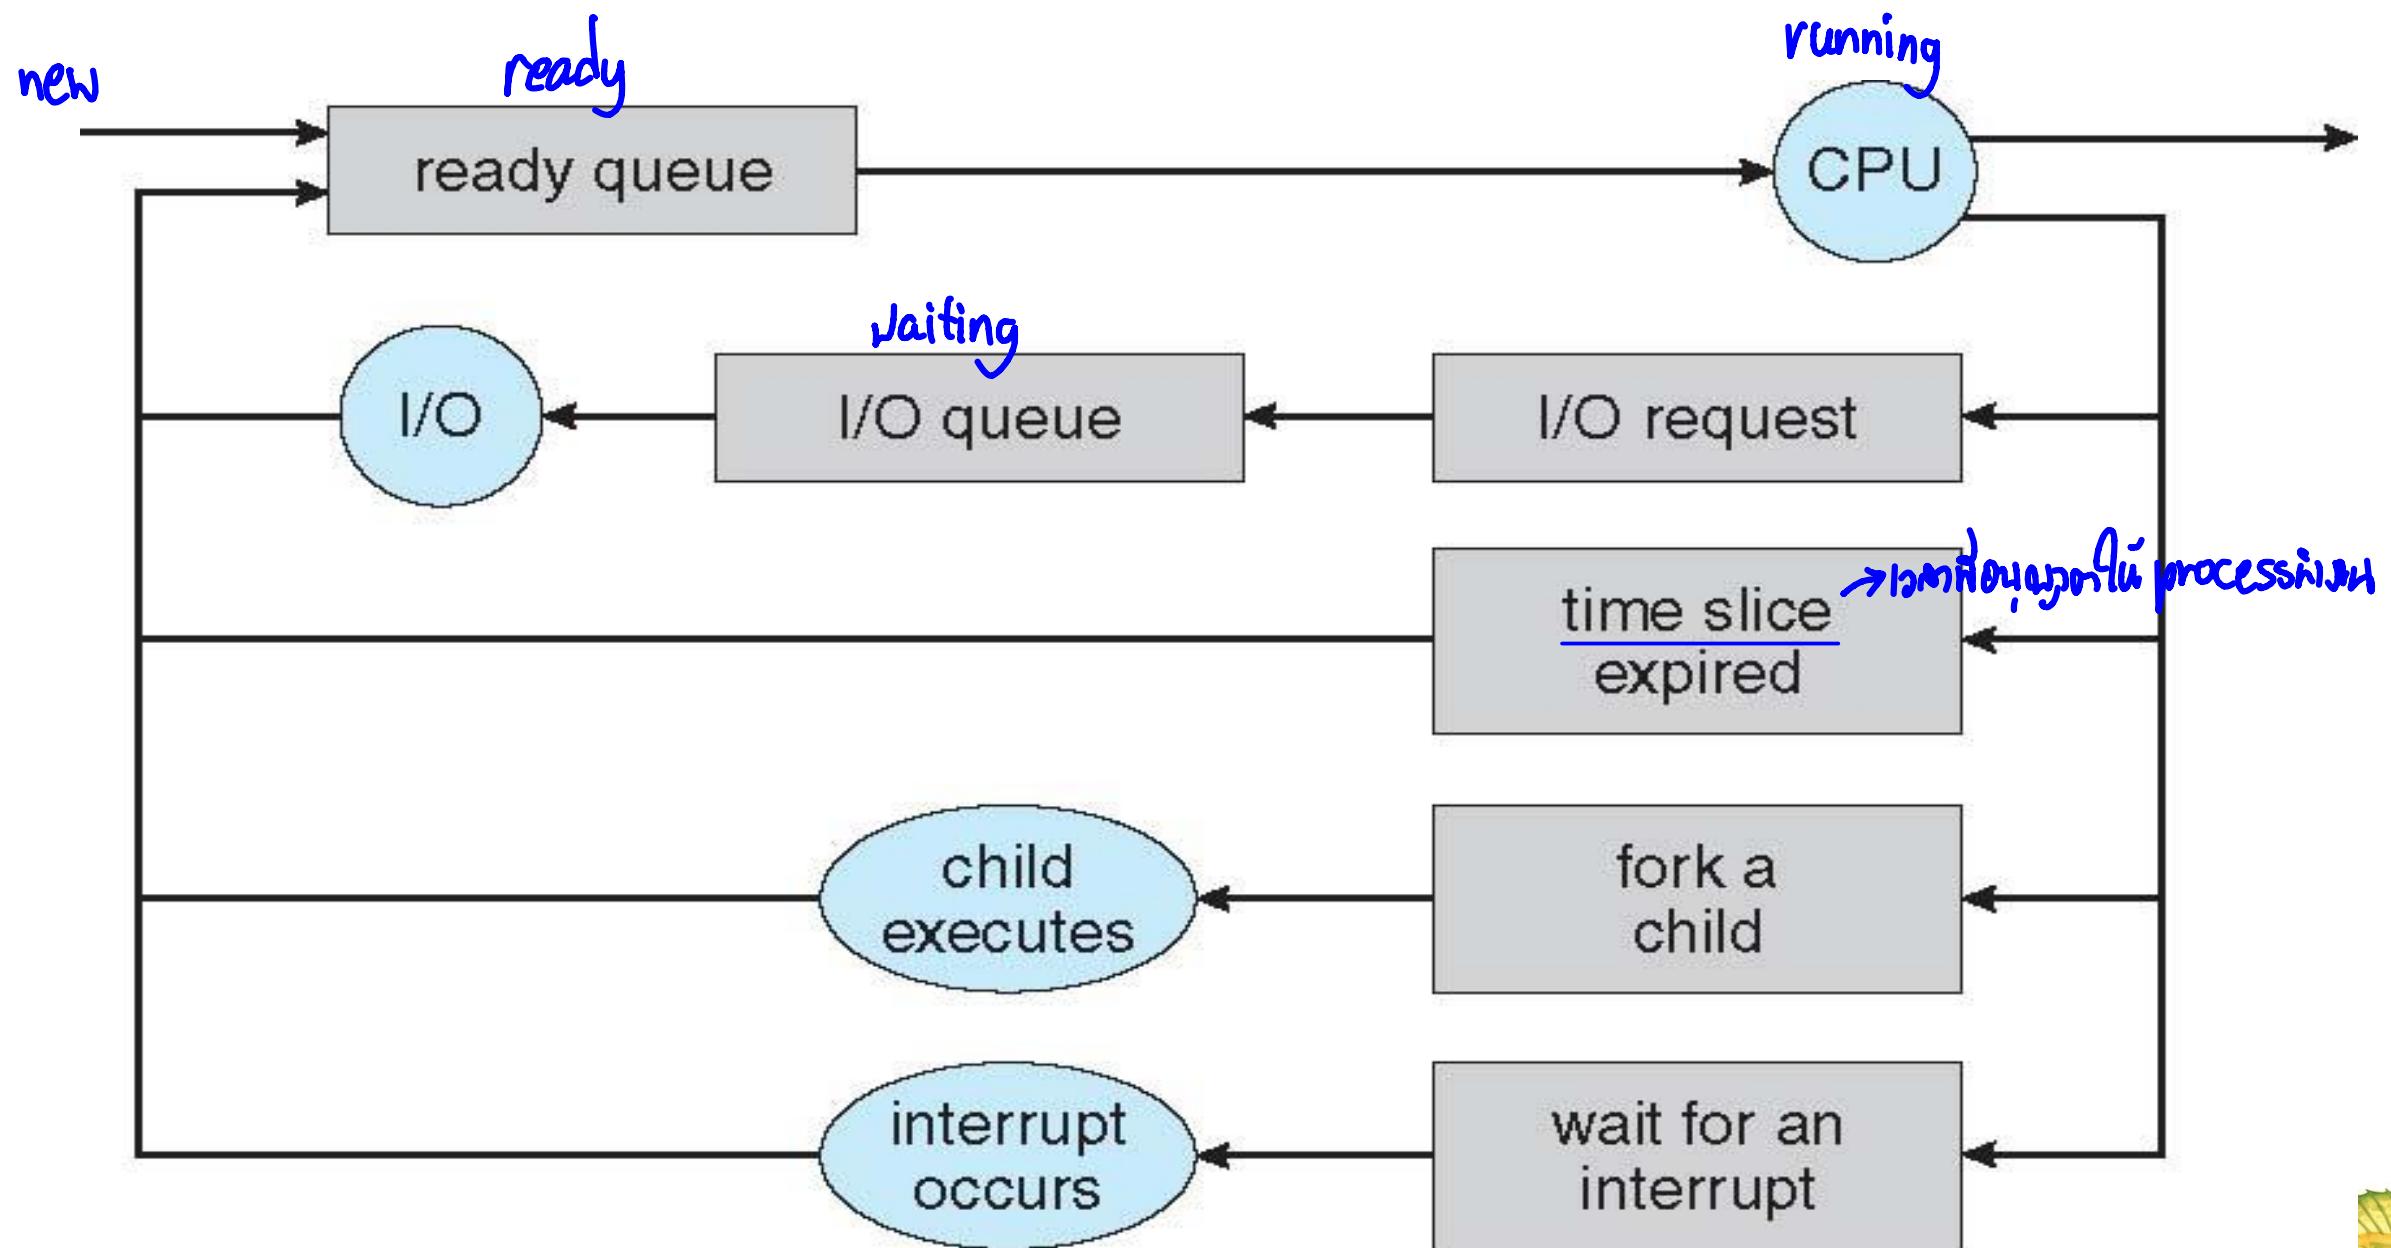
list of PCB





# Representation of Process Scheduling

- Queuing diagram represents queues, resources, flows





# Schedulers

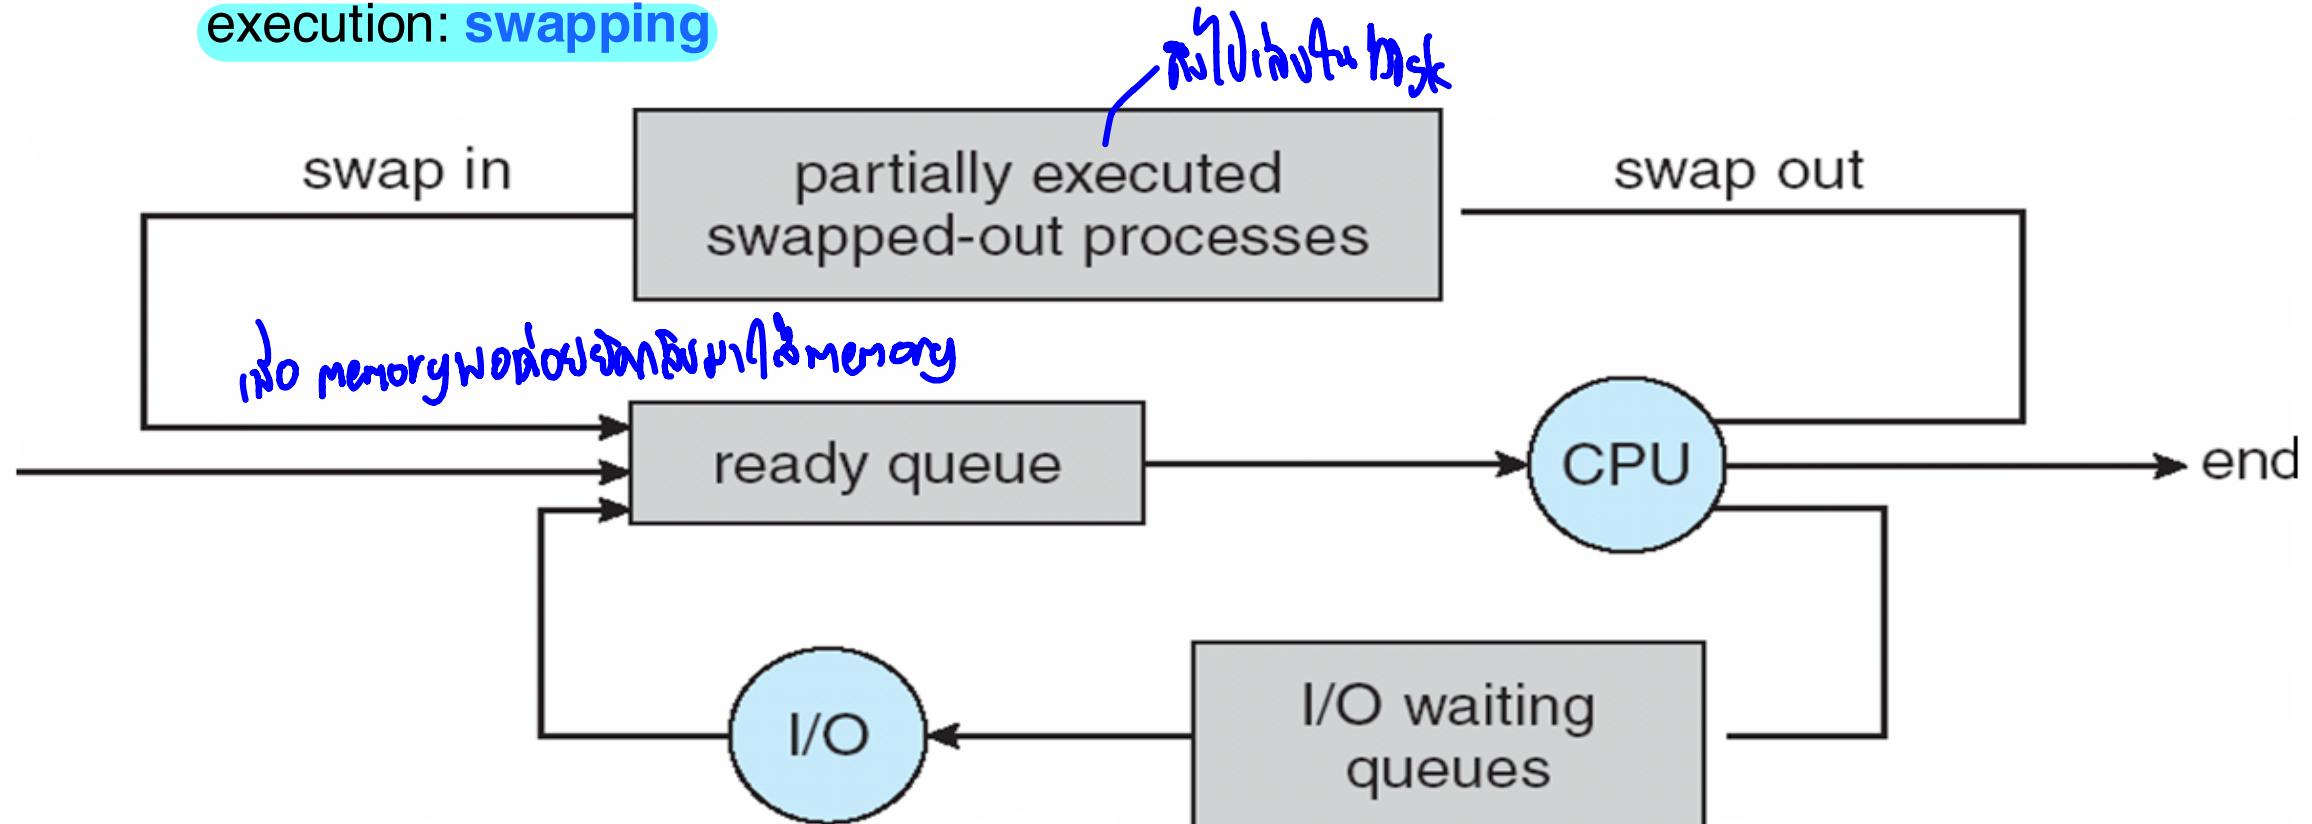
- **Long-term scheduler** (or **job scheduler**) – selects which processes should be brought into the ready queue
- \***Short-term scheduler** (or **CPU scheduler**) – selects which process should be executed next and allocates CPU
  - Sometimes the only scheduler in a system
  - Short-term scheduler is invoked very frequently (milliseconds)  $\Rightarrow$  (must be fast)
- Long-term scheduler is invoked very infrequently (seconds, minutes)  $\Rightarrow$  (may be slow)  
*from Process Management*
- The long-term scheduler controls the **degree of multiprogramming**
- Processes can be described as either:
  - **I/O-bound process** – spends more time doing I/O than computations, many short CPU bursts
  - **CPU-bound process** – spends more time doing computations; few very long CPU bursts
- Long-term scheduler strives for good **process mix**
  - I/O and CPU*





# Addition of Medium Term Scheduling

- **Medium-term scheduler** can be added if degree of multiple programming needs to decrease
  - Remove process from memory, store on disk, bring back in from disk to continue execution: **swapping**





# process/PCB Context Switch

- When CPU switches to another process, the system must **save the state** of the old process and load the **saved state** for the new process via a **context switch**
- **Context** of a process represented in the PCB
  - **Context-switch time** is overhead; the system does no useful work while switching
    - The more complex the OS and the PCB -> longer the context switch
  - Thread-level context switching time:
- Time dependent on hardware support
  - Some hardware provides multiple sets of registers per CPU -> multiple contexts loaded at once

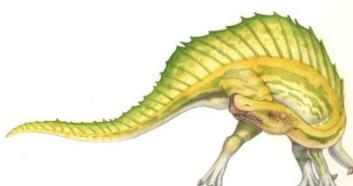




# Operations on Processes

---

- System must provide mechanisms for process creation, termination, and so on as detailed next





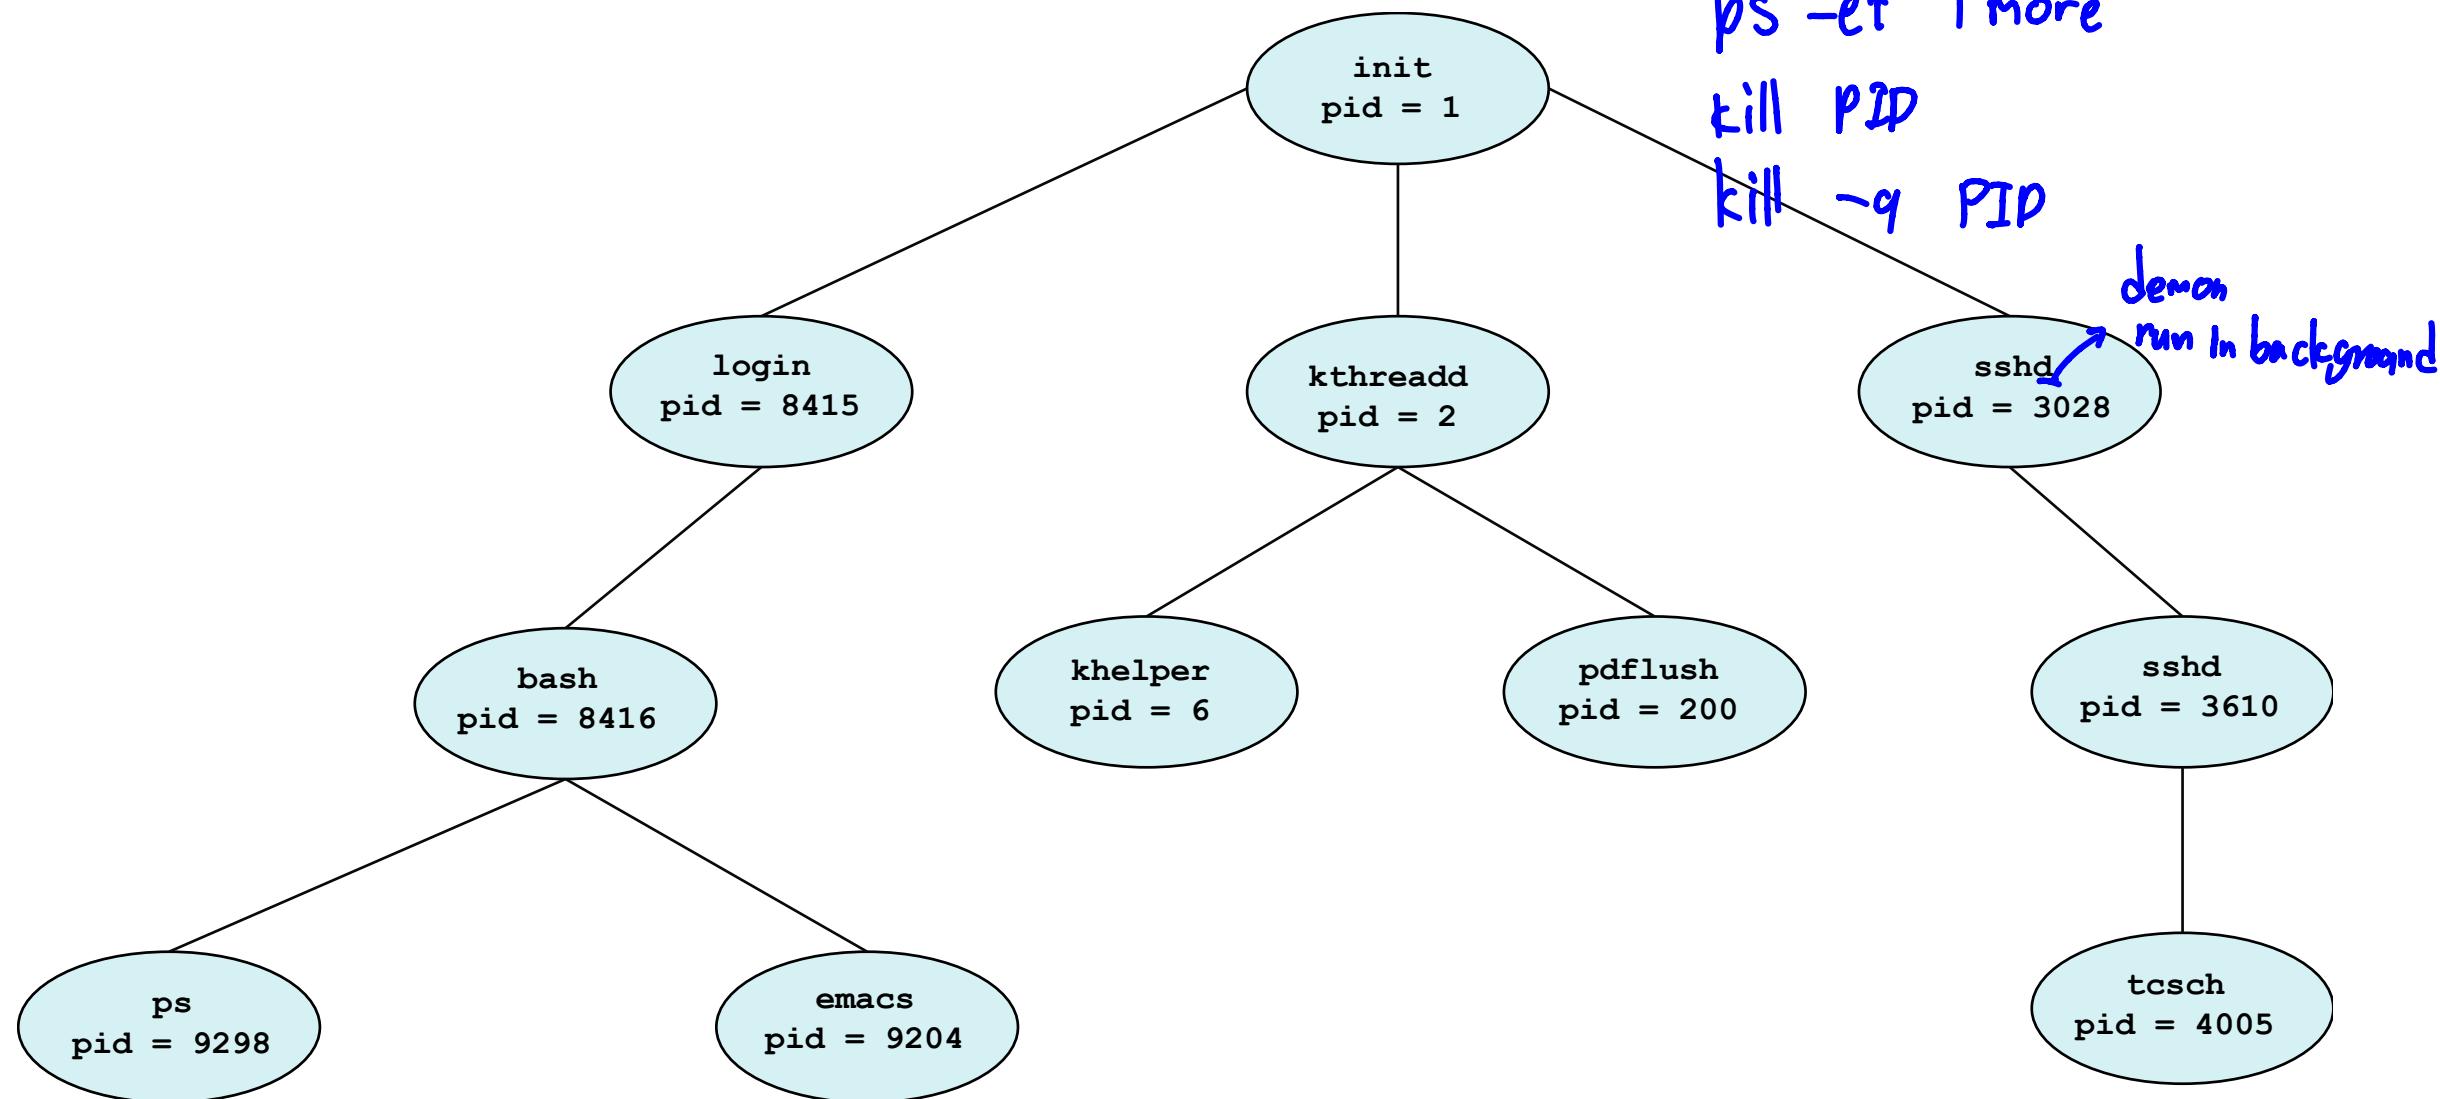
# Process Creation

- Parent process create children processes, which, in turn create other processes, forming a tree of processes
- Generally, process identified and managed via a process identifier (pid)  
mostly int, 16-bit
- Resource sharing options
  - Parent and children share all resources
  - Children share subset of parent's resources
  - Parent and child share no resources
- Execution options
  - Parent and children execute concurrently
  - Parent waits until children terminate





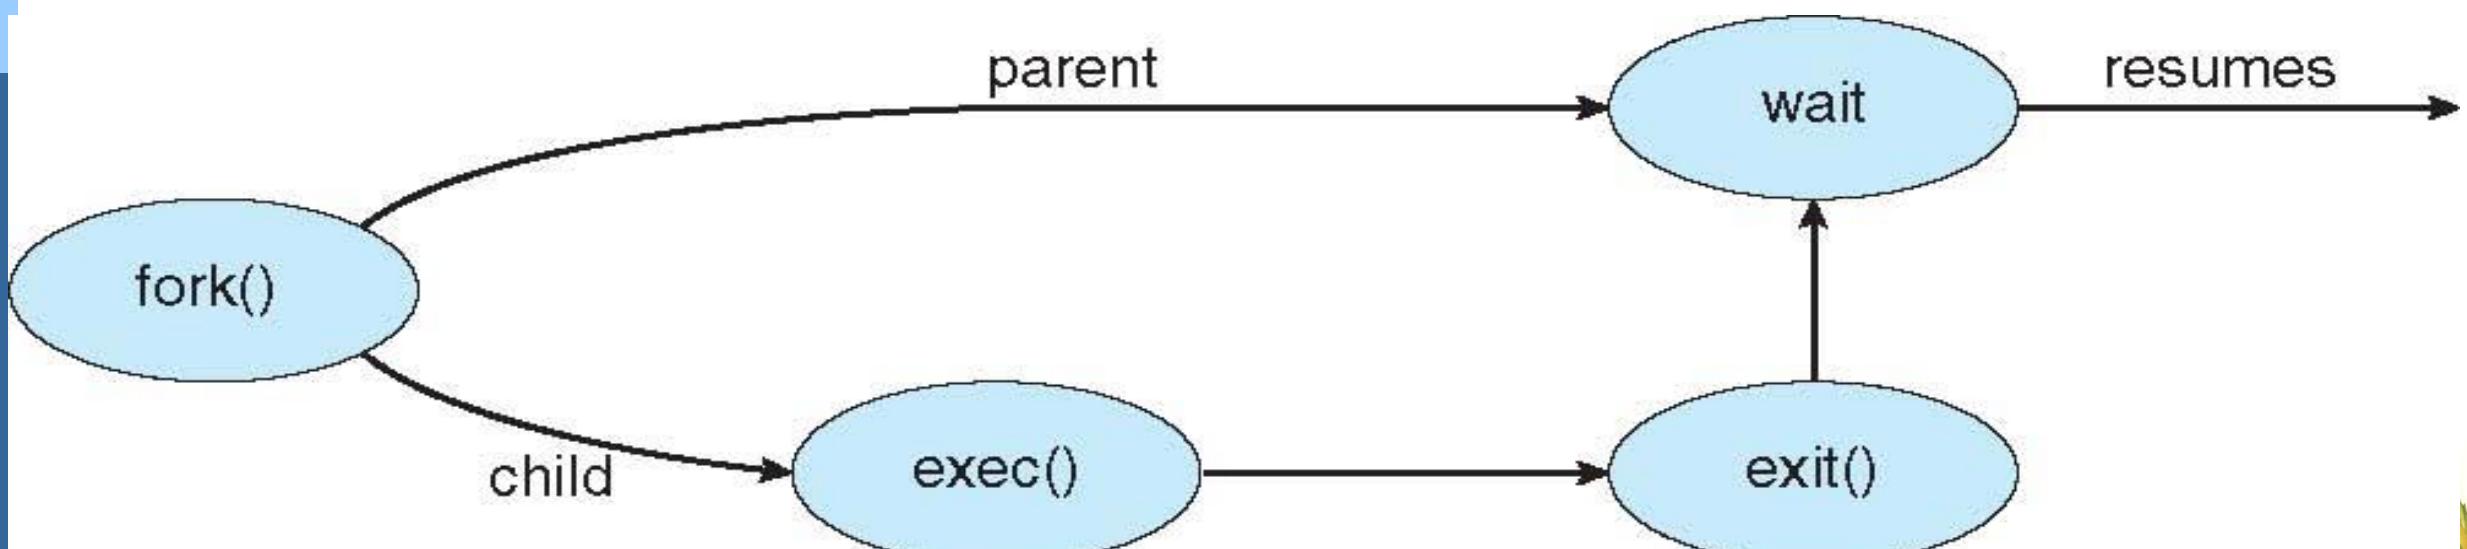
# A Tree of Processes in Linux

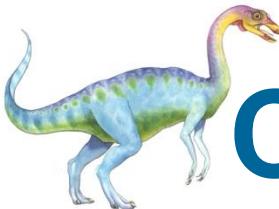




# Process Creation (Cont.)

- Address space
  - Child duplicate of parent
  - Child has a program loaded into it
- UNIX examples
  - `fork()` system call creates new process
  - `exec()` system call used after a `fork()` to replace the process' memory space with a new program





# C Program Forking Separate Process

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

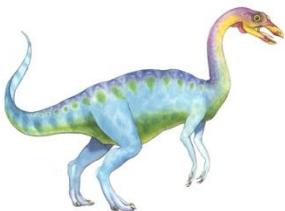
int main()
{
pid_t pid;

/* fork a child process */
pid = fork();

if (pid < 0) { /* error occurred */
    fprintf(stderr, "Fork Failed");
    return 1;
}
else if (pid == 0) { /* child process */
    execlp("/bin/ls", "ls", NULL);
}
else { /* parent process */
    /* parent will wait for the child to complete */
    wait(NULL);
    printf("Child Complete");
}

return 0;
}
```





# Process Termination

- Process executes last statement and asks the operating system to delete it (`exit()`)
  - Output data from child to parent (via `wait()`)
  - Process' resources are deallocated by operating system
- Parent may terminate execution of children processes (`abort()`)
  - Child has exceeded allocated resources
  - Task assigned to child is no longer required
  - If parent is exiting
    - ▶ Some operating systems do not allow child to continue if its parent terminates
      - All children terminated - **cascading termination**

ນິກາງດູໂທ່ານ, ໂມກໍລວມກຳນົດໃນ Process ດັ່ງນີ້ແລ້ວ

- Wait for termination, returning the pid:

```
pid_t pid; int status;  
pid = wait(&status);
```

- If no parent waiting, then terminated process is a **zombie**
- If parent terminated, processes are **orphans**

ເລື່ອດັ່ງນີ້

process) ມີລາຍງາ  
ມີລາຍງາ  
ຫຼັງນີ້ຈະມີລາຍງາ





# Interprocess Communication *IPC*

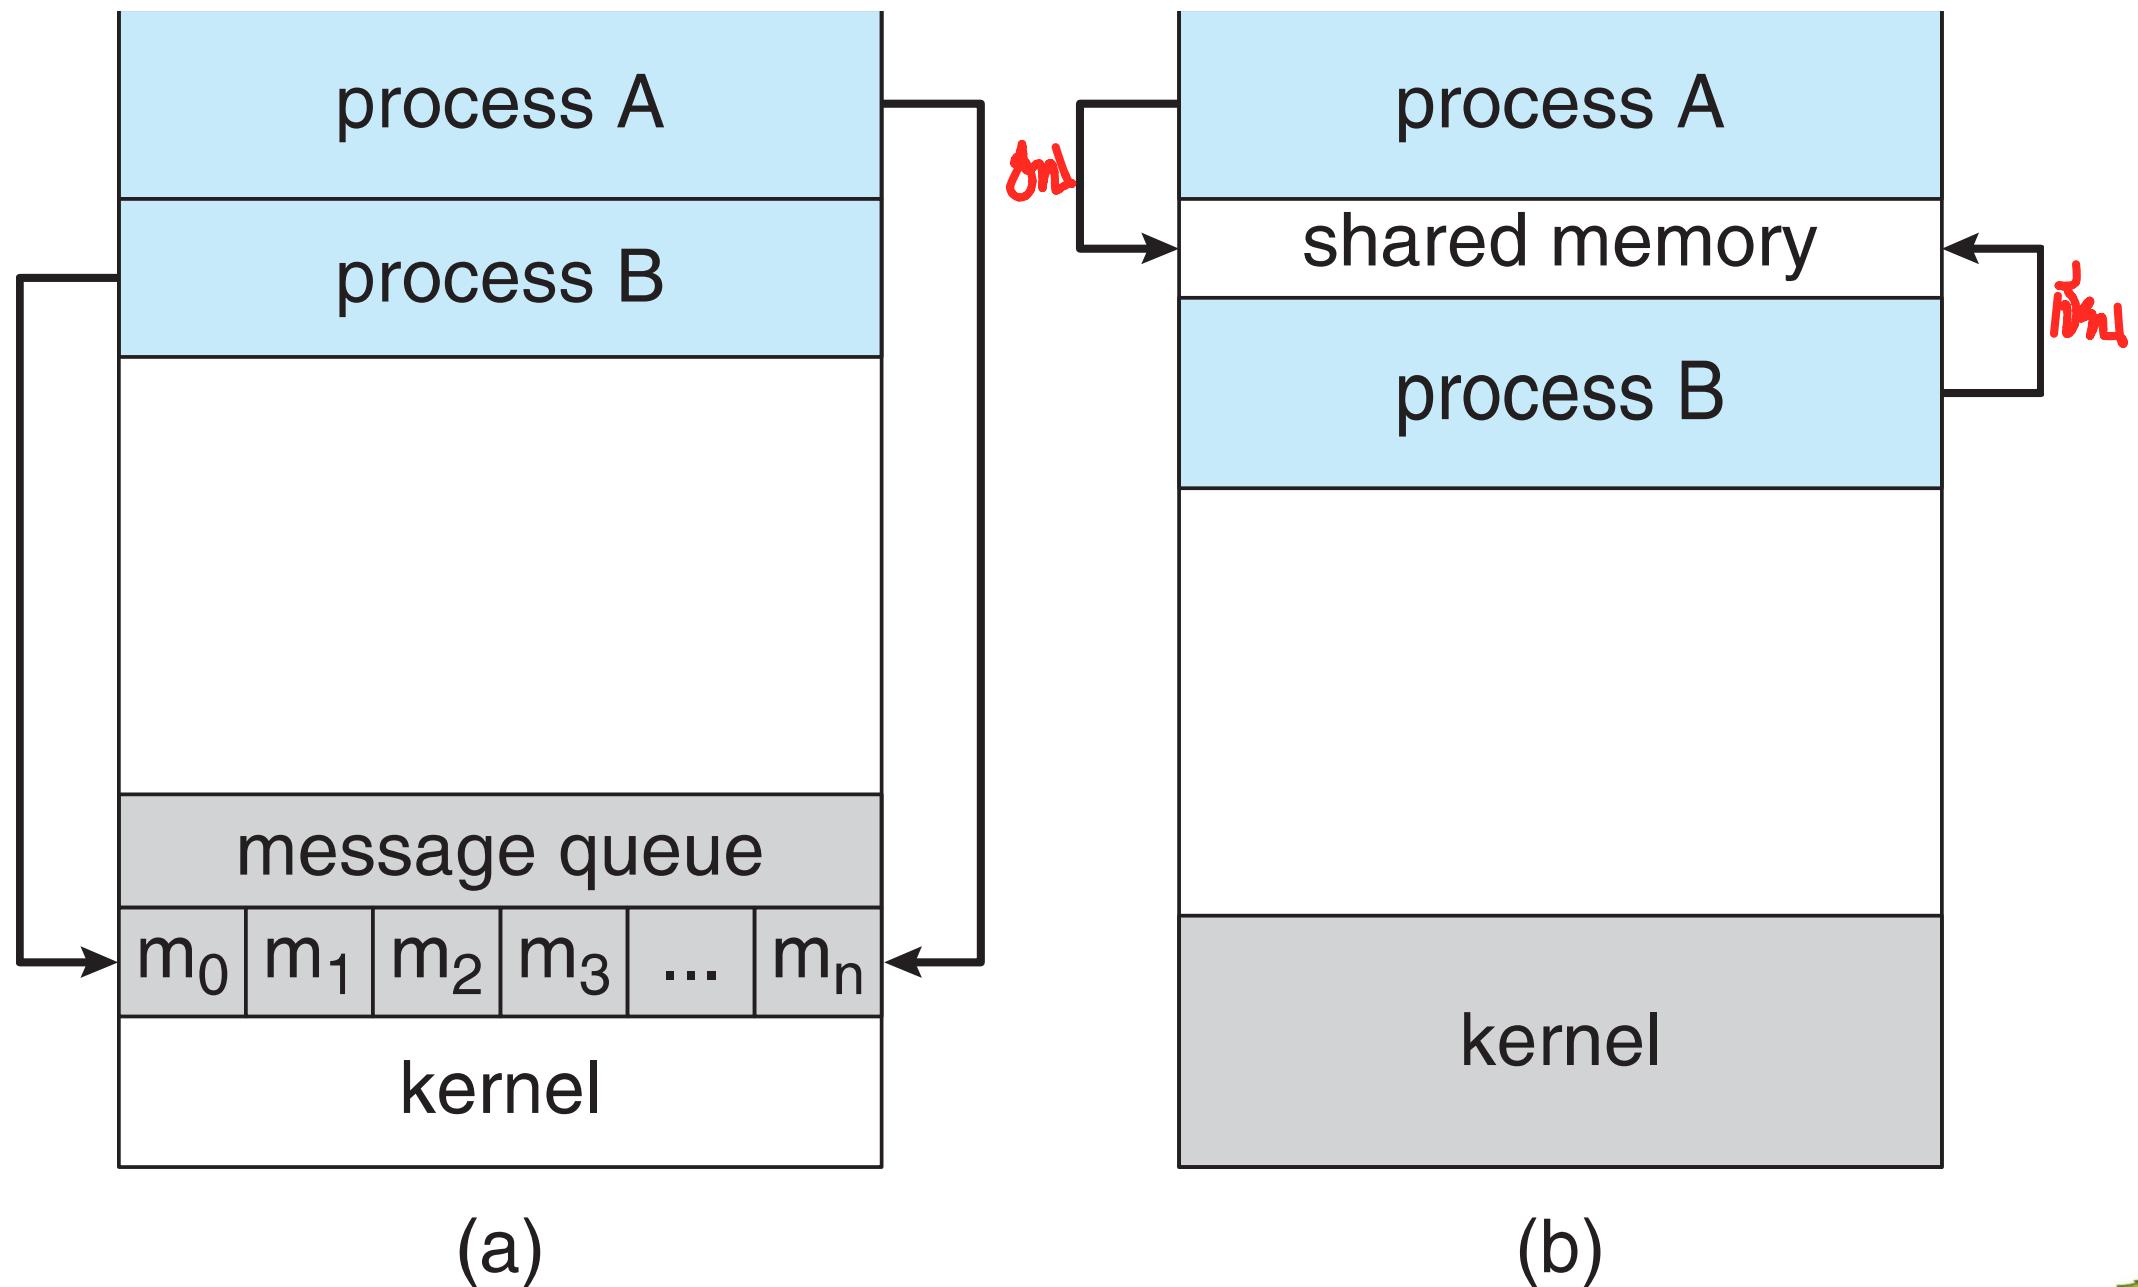
- Processes within a system may be *independent* or *cooperating*
- Cooperating process can affect or be affected by other processes, including sharing data
- Reasons for cooperating processes:
  - Information sharing
  - Computation speedup
  - Modularity
  - Convenience
- Cooperating processes need **interprocess communication (IPC)**
- Two models of IPC ရှုချင်နာရု ၃ မျိုး
  - Shared memory မျတ်စွေးစံဆိုင်
  - Message passing ဓမ္မဂေါ်





# Communications Models

Shared mem



message queue

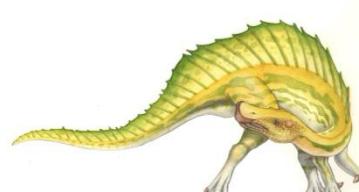
(b)





# Cooperating Processes

- **Independent** process cannot affect or be affected by the execution of another process
- **Cooperating** process can affect or be affected by the execution of another process
- Advantages of process cooperation 
  - Information sharing
  - Computation speed-up
  - Modularity
  - Convenience

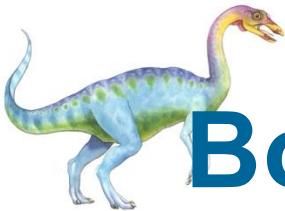




# Producer-Consumer Problem

- Paradigm for cooperating processes, *producer* process produces information that is consumed by a *consumer* process
  - **unbounded-buffer** places **no practical limit** on the size of the buffer
  - **bounded-buffer** assumes that there is a **fixed buffer size**





# Bounded-Buffer – Shared-Memory Solution

- Shared data

```
#define BUFFER_SIZE 10

typedef struct {

    . . .

} item;

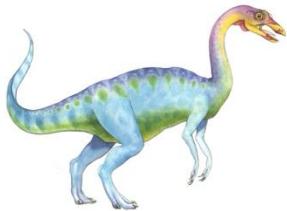
item buffer[BUFFER_SIZE];

int in = 0;

int out = 0;
```

- Solution is correct, but can only use BUFFER\_SIZE-1 elements

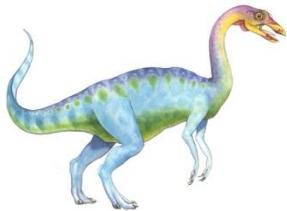




# Bounded-Buffer – Producer

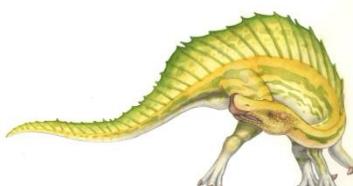
```
item next produced;  
while (true) {  
    /* produce an item in next produced */  
    while (((in + 1) % BUFFER SIZE) == out)  
        ; /* do nothing */  
    buffer[in] = next produced;  
    in = (in + 1) % BUFFER SIZE;  
}
```

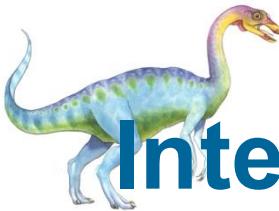




# Bounded Buffer – Consumer

```
item next_consumed;  
  
while (true) {  
    while (in == out)  
        ; /* do nothing */  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
  
    /* consume the item in next_consumed */  
}
```





# Interprocess Communication – Message Passing

- Mechanism for processes to communicate and to synchronize their actions
- Message system – processes communicate with each other without resorting to shared variables
- IPC facility provides two operations:
  - `send(message)` – message size fixed or variable
  - `receive(message)`
- If  $P$  and  $Q$  wish to communicate, they need to:
  - establish a **communication link** between them
  - exchange messages via send/receive

physical, TCP protocol
- Implementation of communication link
  - physical (e.g., shared memory, hardware bus)
  - logical (e.g., direct or indirect, synchronous or asynchronous, automatic or explicit buffering)





# Implementation Questions

- How are links established?
- Can a link be associated with more than two processes?
- How many links can there be between every pair of communicating processes?
- What is the capacity of a link?
- Is the size of a message that the link can accommodate fixed or variable?
- Is a link unidirectional or bi-directional?





# Direct Communication

↑  
Tnjōwñ

↑  
Baoñ

- Processes must name each other explicitly:
  - **send(P, message)** – send a message to process P
  - **receive(Q, message)** – receive a message from process Q
  
- Properties of communication link
  - Links are established automatically
  - A link is associated with exactly one pair of communicating processes
  - Between each pair there exists exactly one link
  - The link may be unidirectional, but is usually bi-directional

↑  
Iny

↑  
Injōwñ



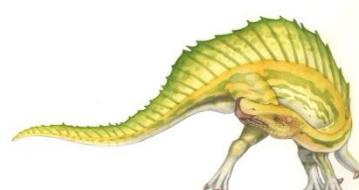


# Indirect Communication

- Messages are directed and received from mailboxes (also referred to as ports)
  - Each mailbox has a unique id
  - Processes can communicate only if they share a mailbox
  
- Properties of communication link
  - Link established only if processes share a common mailbox
  - A link may be associated with many processes
  - Each pair of processes may share several communication links
  - Link may be unidirectional or bi-directional

ଟ୍ରେକିଙ୍

socket





# Indirect Communication

↳ *Indirect communication*

*within shared memory area, shared memory within itself*

- Operations

- create a new mailbox
  - send and receive messages through mailbox
  - destroy a mailbox

- Primitives are defined as:

**send(A, message)** – send a message to mailbox A

**receive(A, message)** – receive a message from mailbox A

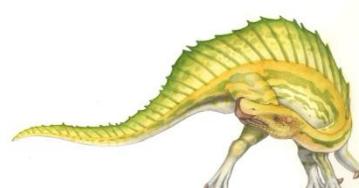




# Indirect Communication

---

- Mailbox sharing
  - $P_1$ ,  $P_2$ , and  $P_3$  share mailbox A
  - $P_1$ , sends;  $P_2$  and  $P_3$  receive
  - Who gets the message?
  
- Solutions
  - Allow a link to be associated with at most two processes
  - Allow only one process at a time to execute a receive operation
  - Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was.





# Synchronization

- Message passing may be either blocking or non-blocking
- **Blocking** is considered **synchronous**
  - **Blocking send** has the sender block until the message is received
  - **Blocking receive** has the receiver block until a message is available
- **Non-blocking** is considered **asynchronous**
  - **Non-blocking** send has the sender send the message and continue
  - **Non-blocking** receive has the receiver receive a valid message or null

}





# Synchronization (Cont.)

- Different combinations possible
  - If both send and receive are blocking, we have a **rendezvous**
- Producer-consumer becomes trivial

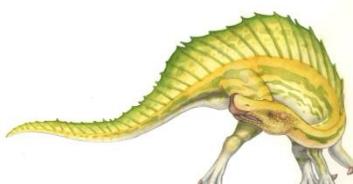
```
message next produced;  
while (true) {  
    /* produce an item in next produced */  
    send(next produced);  
}  
  
message next consumed;  
while (true) {  
    receive(next consumed);  
  
    /* consume the item in next consumed */  
}
```

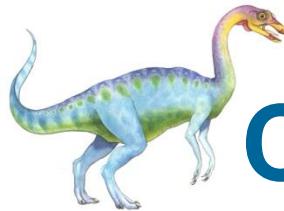




# Buffering

- Queue of messages attached to the link;  
implemented in one of three ways
  1. Zero capacity – 0 messages  
Sender must wait for receiver (rendezvous)
  2. Bounded capacity – finite length of  $n$  messages  
Sender must wait if link full
  3. Unbounded capacity – infinite length  
Sender never waits





# Communications in Client-Server Systems

- Remote Procedure Calls
- Pipes



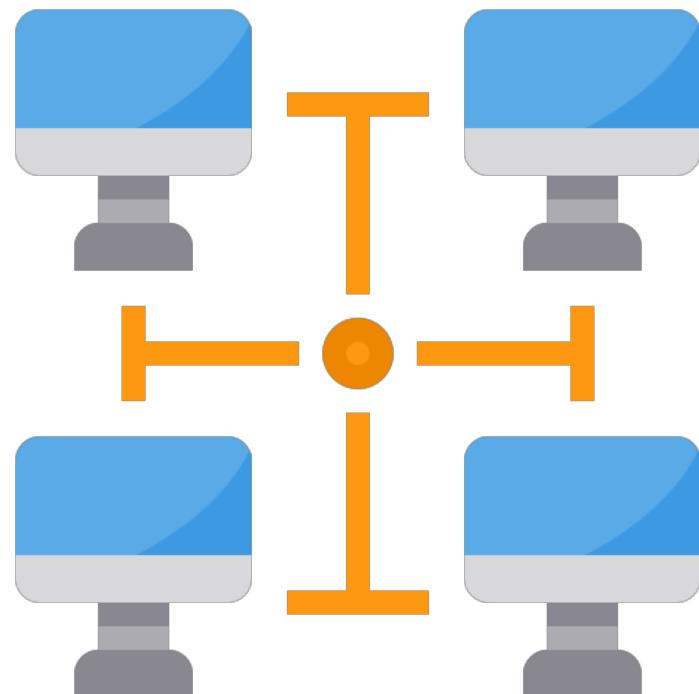


# Remote Procedure Calls

- Remote procedure call (**RPC**) abstracts procedure calls between processes on networked systems
  - Again uses ports for service differentiation
- **Stubs** – client-side proxy for the actual procedure on the server
- The client-side stub locates the server and **marshalls** the parameters
- The server-side stub receives this message, unpacks the marshalled parameters, and performs the procedure on the server
- On Windows, stub code compile from specification written in **Microsoft Interface Definition Language (MIDL)**
- Data representation handled via **External Data Representation (XDL)** format to account for different architectures
  - **Big-endian** and **little-endian**
- Remote communication has more failure scenarios than local
  - Messages can be delivered **exactly once** rather than **at most once**
- OS typically provides a rendezvous (or **matchmaker**) service to connect client and server



# Remote Procedure Call (RPC)

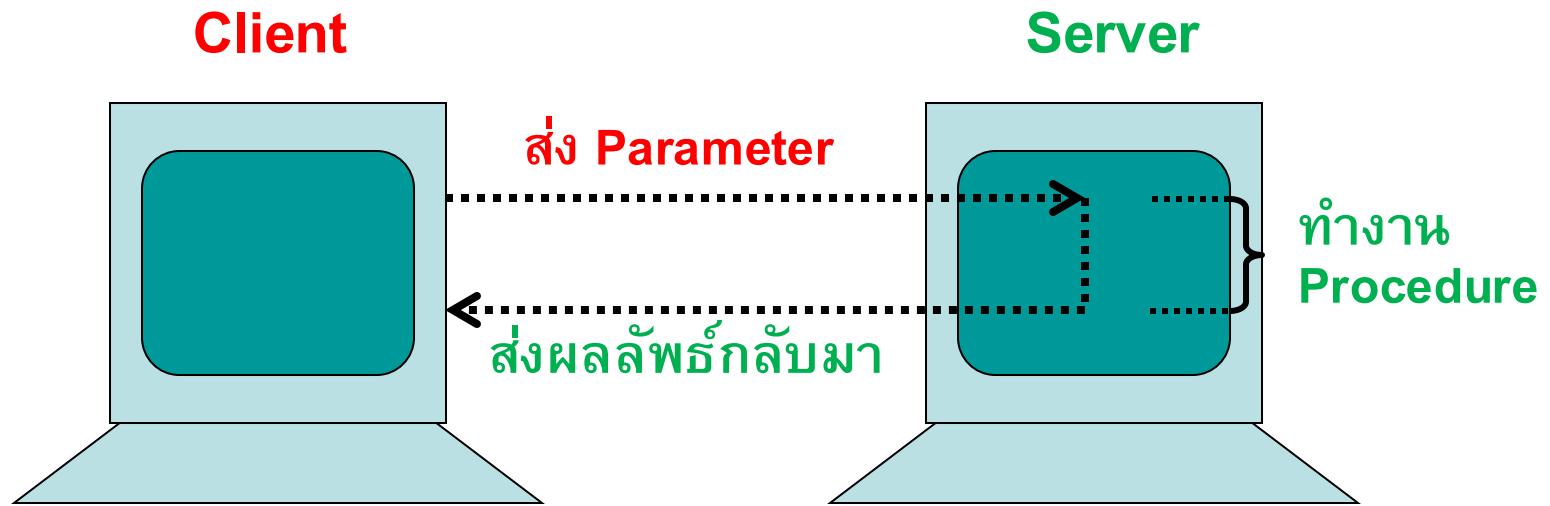


# Motivation for RPC

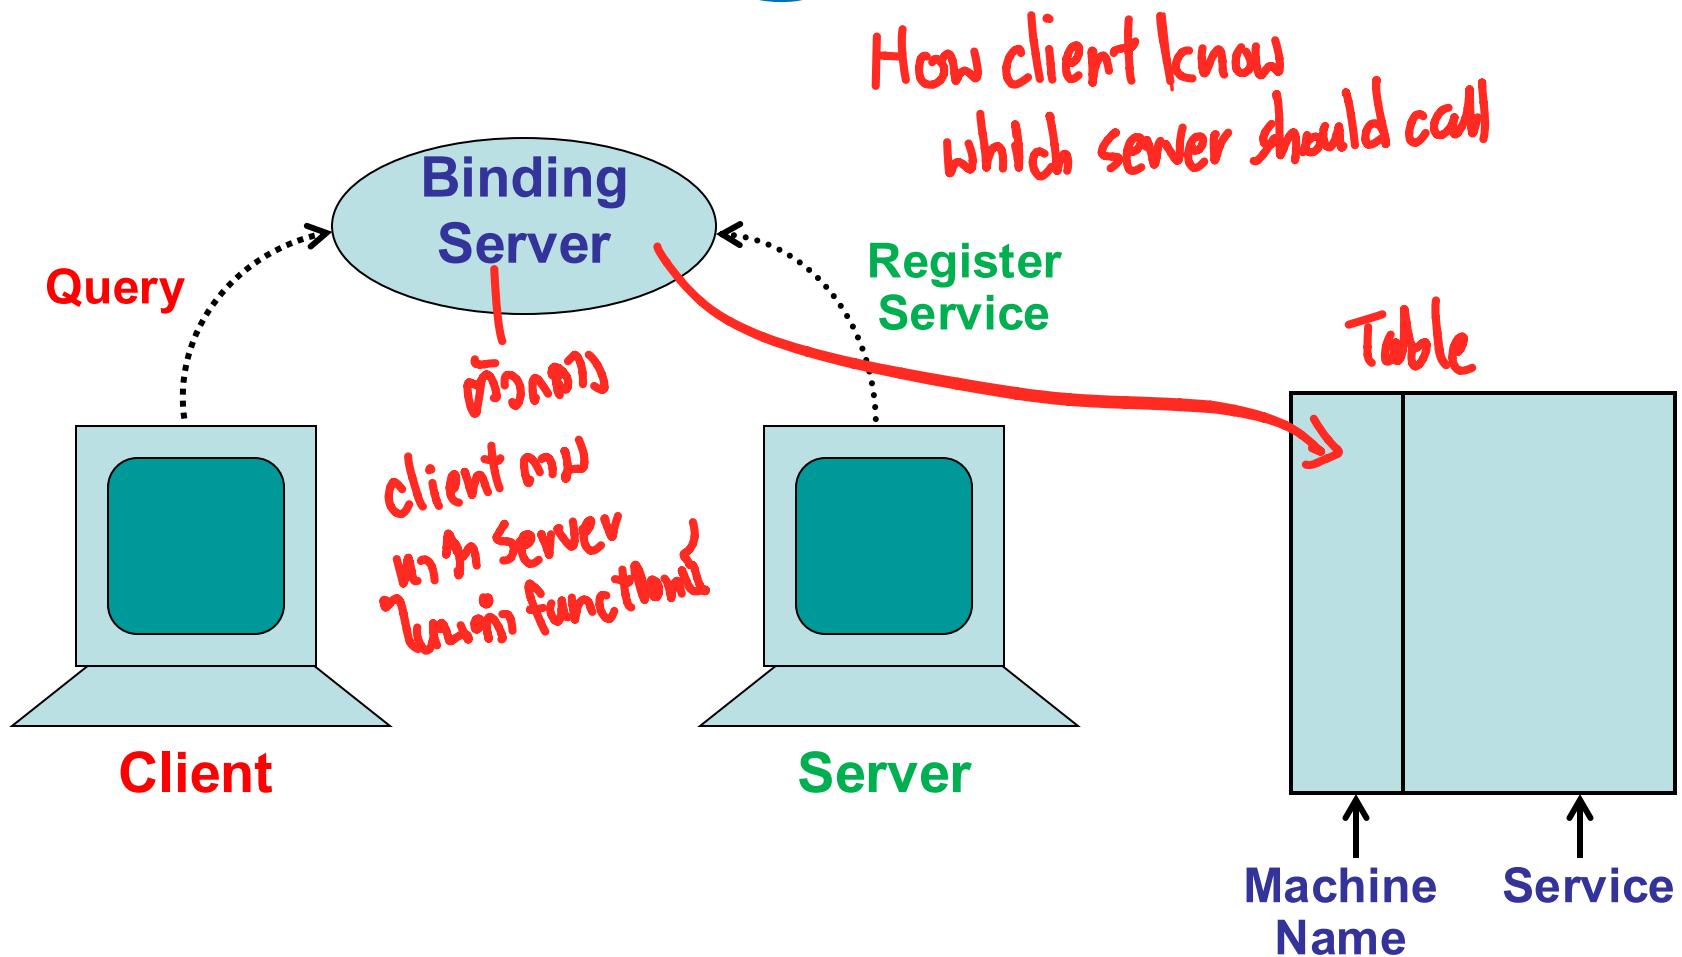
function call

- Procedure call เป็นวิธีที่รู้จักกันดีอยู่แล้ว **ในการ transfer control และ data ในคอมพิวเตอร์เครื่องเดียวกัน**
- ดังนั้น **ขยายแนวคิดนี้ไปใน network** **ต่อไป**

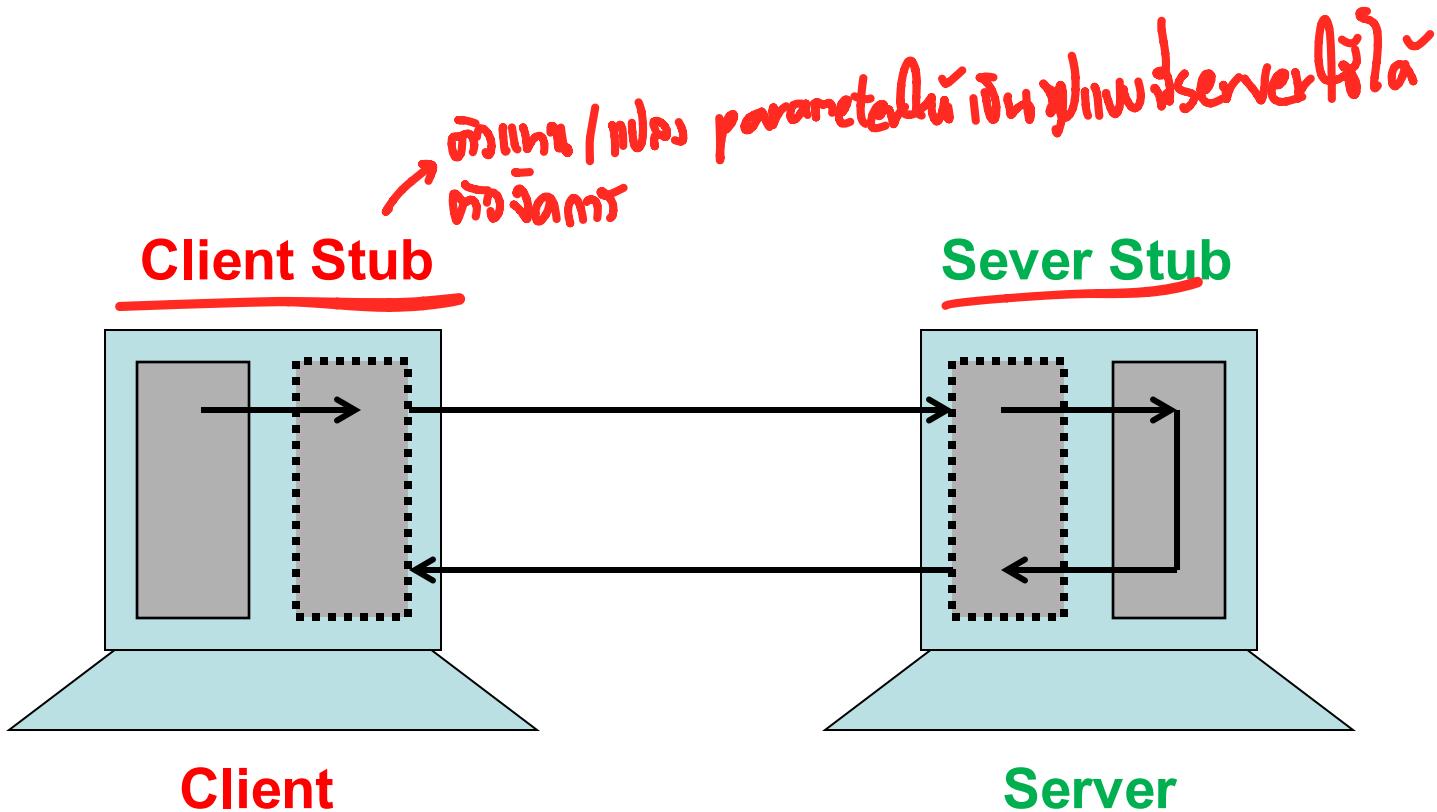
# Basic RPC Operation



# Binding Server



# การใช้ stub



# Stub procedure

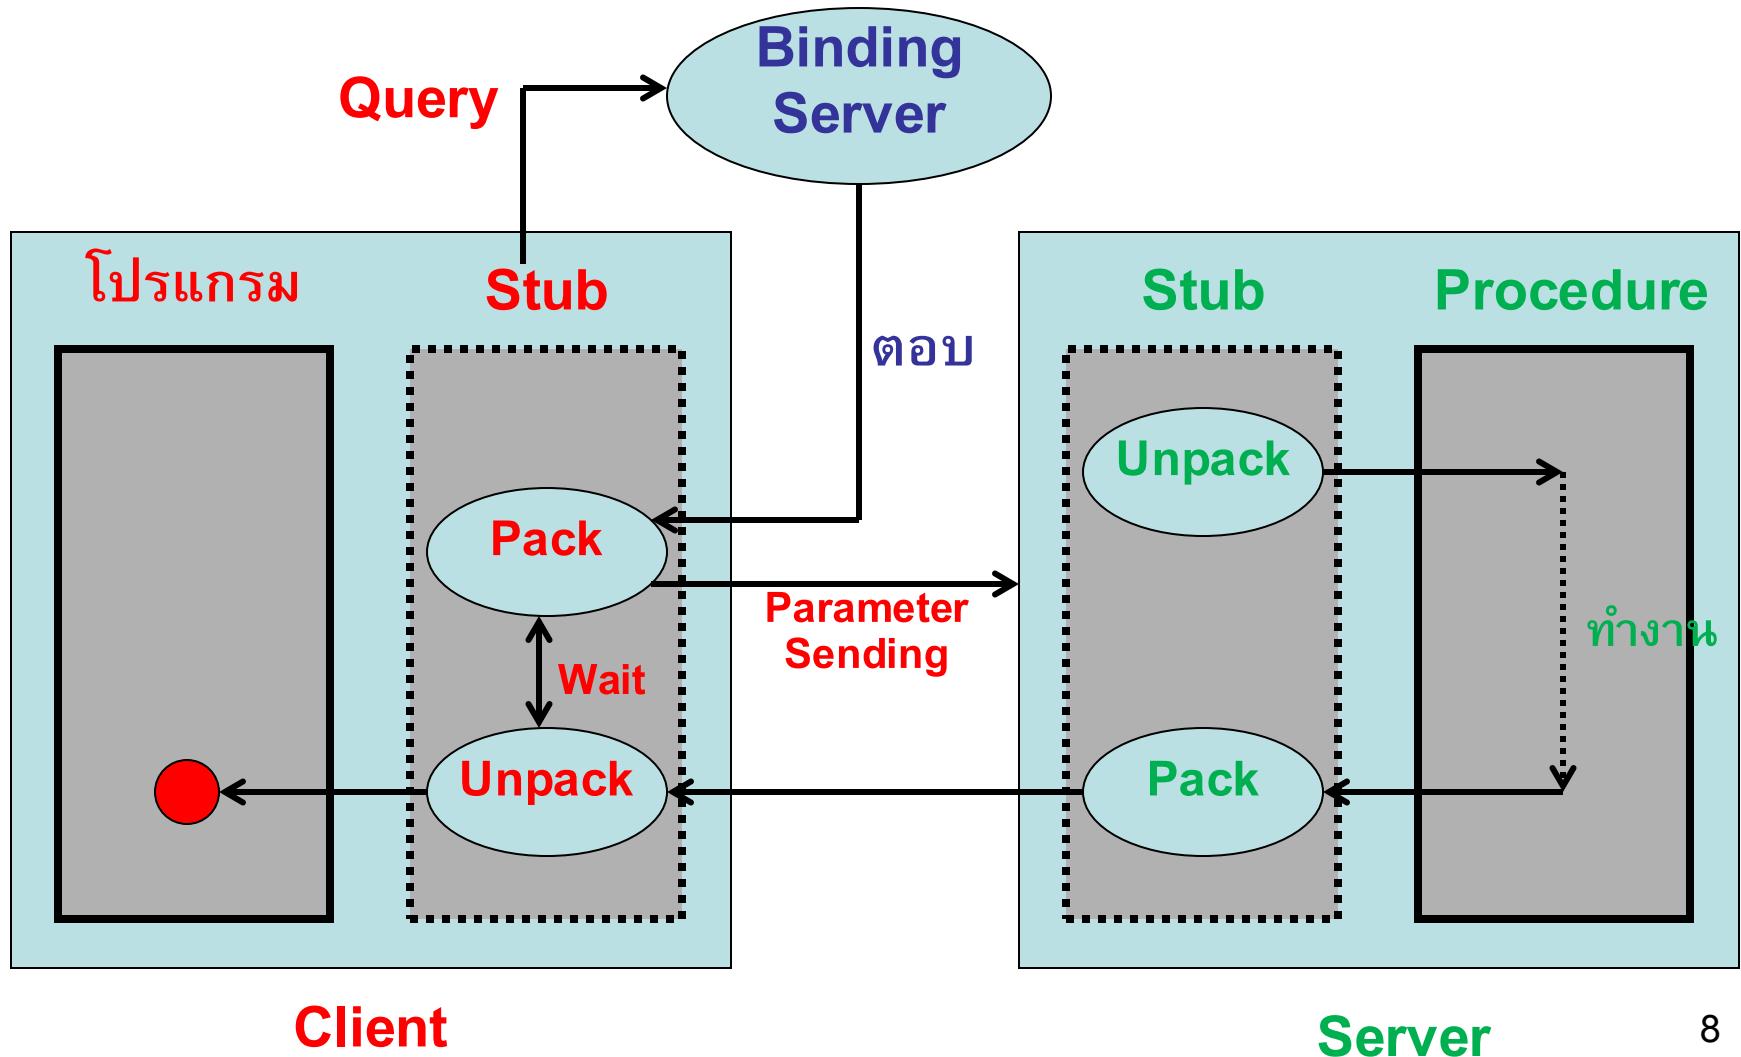
- **Dummy procedure** ใน client และ server
- ทำหน้าที่สร้าง message ประกอบด้วย parameter และส่งให้ server
- Server stub ส่ง parameter ให้ procedure จริงทำงานและส่งผลลัพธ์กลับให้ client

# การทำ parameter passing

- แปลงเป็น network representation (XDR) ที่เข้าใจกันทั่วสองฝ่าย
- Pack parameter ให้สามารถส่งใน network ได้อย่างเหมาะสม เรียกว่า parameter marshalling

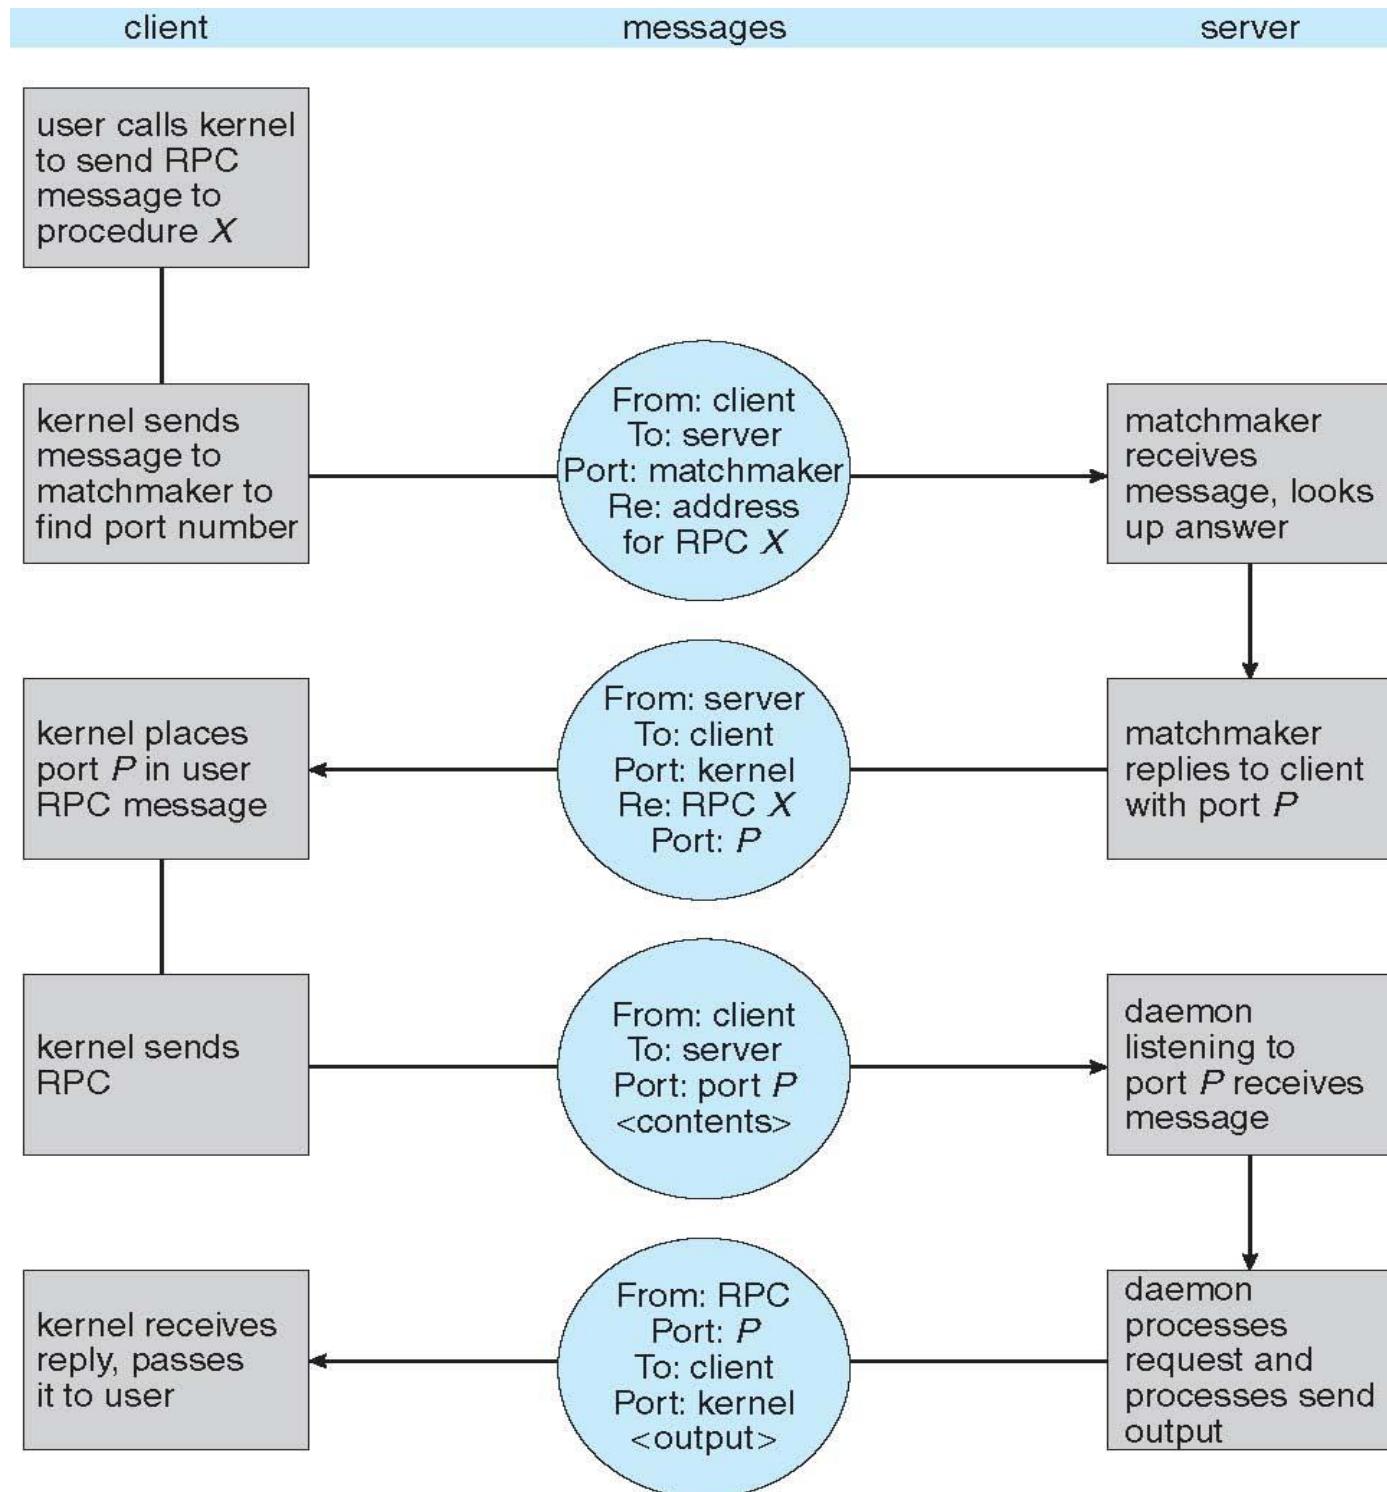
หน้า stub

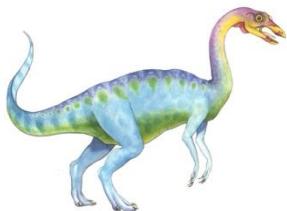
# การทำงานของ RPC





# Execution of RPC





# Pipes

గ్రంథాలు

head -3 filepath | tail -1

cat | head -1

- Acts as a conduit allowing two processes to communicate

ప్రాణీకరణ

## □ Issues

UNIX

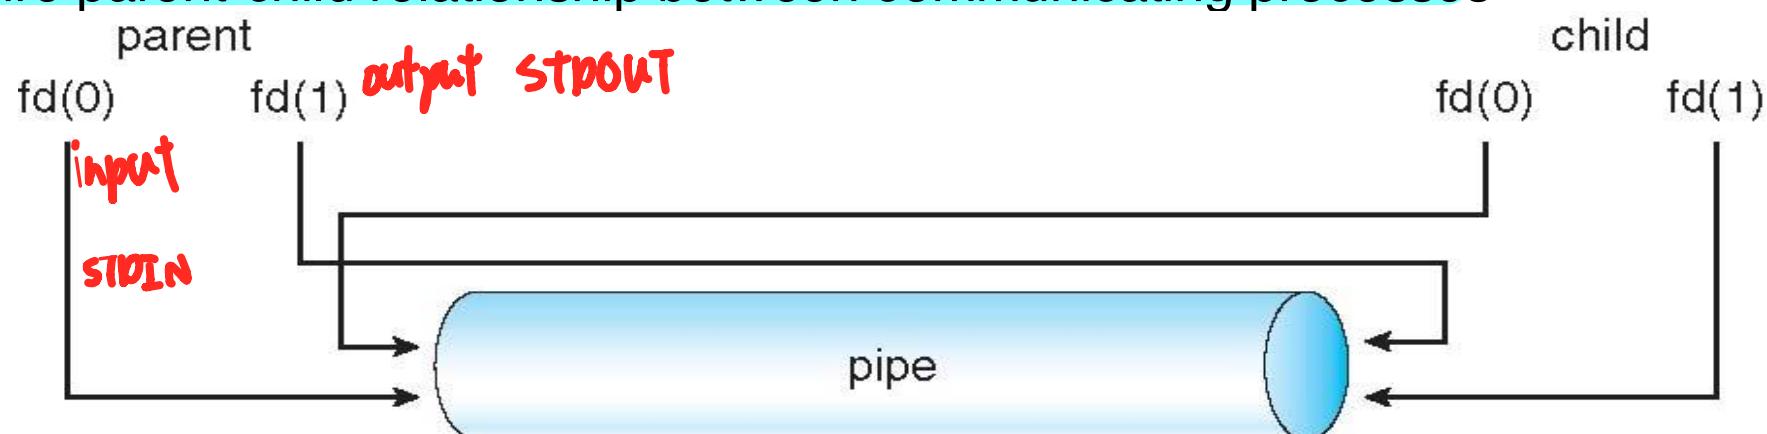
- Is communication unidirectional or bidirectional?
- In the case of two-way communication, is it half or full-duplex?
- Must there exist a relationship (i.e. **parent-child**) between the communicating processes? **No**
- Can the pipes be used over a network? **No**



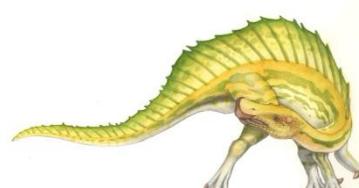


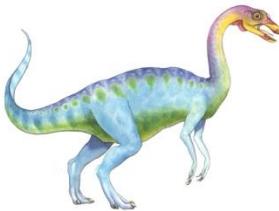
# Ordinary Pipes

- Ordinary Pipes allow communication in standard producer-consumer style
- Producer writes to one end (the **write-end** of the pipe)
- Consumer reads from the other end (the **read-end** of the pipe)
- Ordinary pipes are therefore unidirectional
- Require parent-child relationship between communicating processes



- Windows calls these **anonymous pipes**
- See Unix and Windows code samples in textbook





# Named Pipes

- Named Pipes are more powerful than ordinary pipes
- Communication is bidirectional
- No parent-child relationship is necessary between the communicating processes
- Several processes can use the named pipe for communication
- Provided on both UNIX and Windows systems

